

**DOOMSCROLLING DETECTION AND DIGITAL MINDFULNESS
MOBILE APPLICATION FOR SHORT-FORM VIDEO PLATFORMS
USING VADER AND FUZZY LOGIC**

by

Cadigal, Genesis A.
Chuayap, Rayan Kennard O.
Limos, Luigi Karl B.
Mapesos, Gean Dhylan E.

Submitted in Partial Fulfillment of the Requirements for the Degree of
Bachelor of Science in Computer Science with Specialization in Software Engineering

at

FEU Institute of Technology

Month/Year

Mrs. Rain Marilyn M. Sanchez
Thesis Adviser

©2026 Cadigal, Genesis A., Chuayap, Rayan Kennard O., Limos, Luigi Karl B. and
Mapesos, Gean Dhylan E.
All Rights Reserved

The authors grant FEU Institute of Technology permission to reproduce and distribute the
contents of this document in whole or in part.

APPROVAL AND ACCEPTANCE SHEET

The thesis entitled “**DOOMSCROLLING DETECTION AND DIGITAL MINDFULNESS MOBILE APPLICATION FOR SHORT-FORM VIDEO PLATFORMS USING VADER AND FUZZY LOGIC**” prepared and submitted by:

Cadigal, Genesis A.
Chuayap, Rayan Kennard O.
Limos, Luigi Karl B.
Mapesos, Gean Dhylan E.

In partial fulfillment of the course of requirement for the **Degree of Bachelor of Science in Computer Science with specialization in Software Engineering** has been examined and is hereby recommended for approval.

MS. MAY FLORENCE SAN PABLO

Panelist 1

PROF. JUSTINE PURA

Panelist 2

MS. JUAN C. DELA CRUZ

Head Panelist

Accepted as partial fulfillment of the requirements for the Degree of Bachelor of Science in Computer Science with specialization in Software Engineering.

MRS. RAIN MARILYN M. SANCHEZ

Thesis Adviser

MRS. ELISA V. MALASAGA

Course Adviser

DR. SHANETH C. AMBAT

Department Head

Date

ACKNOWLEDGMENT

[Acknowledgment text goes here.]

TABLE OF CONTENTS

THESIS TITLE	i
APPROVAL AND ACCEPTANCE SHEET	iii
ACKNOWLEDGMENT	iv
TABLE OF CONTENTS	v
LIST OF TABLES	vii
LIST OF FIGURES	ix
1. INTRODUCTION	1
1.1. Background of the Study	2
1.2. Statement of the Problem	3
1.3. Significance of the Study	4
1.4. Objectives of the Study	5
1.5. Scope and Delimitations	6
1.6. Limitations	8
1.7. Conceptual Framework	11
1.8. Theoretical Framework	12
1.9. Definition of Terms	13
2. REVIEW OF RELATED LITERATURE AND STUDIES	17
2.1. Foreign Literatures	17
2.2. Local Literatures	19
2.3. Interim Synthesis of Literature	20
2.4. Foreign Studies	20
2.5. Local Studies	25
2.6. Integrated Synthesis and Research Gap	27
3. METHODOLOGY	31
3.1. Type of Research	31
3.2. Research Design	31

3.3.	System Architecture	32
3.4.	Hardware and Software Specifications	42
3.5.	Method in Developing Software	43
3.6.	Method in Evaluation of the Study	53
3.7.	Data Collection Scheduling	56
3.8.	Data Gathering	57
3.9.	Respondents of the Study and Sampling Technique	62
3.10.	Statistical Treatment	64
3.11.	Ethical Considerations	69
4.	PRESENTATION, ANALYSIS AND INTERPRETATION OF DATA	71
4.1.	Presentation of Results	71
4.2.	Narrative Interpretation of Results	84
4.3.	Summary of Findings	85
5.	CONCLUSION	87
5.1.	Summary of Findings	87
5.2.	Conclusions Based on Research Objectives and Questions	88
6.	RECOMMENDATIONS	91
6.1.	Recommendations for System Improvement	91
6.2.	Recommendations for Future Development	92
6.3.	Recommendations for Future Researchers	92
6.4.	Recommendations for Future Implementation and Deployment	93
	REFERENCES	96
	APPENDICES	110

LIST OF TABLES

Table 1	Comparative Analysis of Digital Well-Being Tools and the Proposed System	28
Table 2	Design Justification Matrix: Theoretical Constructs to System Variables	39
Table 3	Fuzzy Set Membership Boundaries for Analytic Scoring	49
Table 4	Representative Appendix A Rules (6 of 27)	52
Table 5	Variable–Instrument Alignment Matrix	62
Table 6	System Architecture and Estimation Framework Result Matrix	72
Table 7	Black Box and White Box Testing Summary	73
Table 8	Evaluation Method, Respondents, and Instrument Status	74
Table 9	Scoring Method and Interpretation Rules	74
Table 10	Participant Profile Summary	75
Table 11	Target Platform Use and Extraction Coverage	75
Table 12	Primary Between-Group Change-Score Comparison	76
Table 13	Secondary Within-Intervention Pre/Post Comparison	77
Table 14	Supplementary Prompt-Excluded Active-Use Metrics	77
Table 15	Prompt-Response Summary for Intervention Participants	78
Table 16	Reliability Event Summary	78
Table 17	DSI Coverage and Sentiment-Reliable Session Summary	79
Table 18	Baseline Component and Composite Correlation with Doomscrolling Scale	79
Table 19	ISO/IEC 25010 User Evaluation Summary	80
Table 20	System Usability Scale Result	80
Table 21	Technology Acceptance Model User Evaluation Summary	80
Table 22	Cronbach’s Alpha Results for User Survey Subscales	81
Table 23	Algorithm Complexity Analysis	81
Table 24	Subject Matter Expert Evaluation Summary	82
Table 25	Open-Ended Feedback Coding Summary	83
Table 26	Overall Results Summary	83
Table 27	Summary of Findings by Research Question	85

Table 28 Summary of Findings	87
Table 29 System Improvement Recommendations	91
Table 30 Future Development Recommendations	92
Table 31 Future Research Recommendations	93
Table 32 Future Implementation and Deployment Recommendations	93
Table 33 Full 27-Rule Fuzzy Inference Base	110
Table 34 Fallback 9-Rule Base (Sentiment-Unreliable Sessions)	113

LIST OF FIGURES

Figure 1 Conceptual Framework of the Study using IPO Model	11
Figure 2 Doomscrolling Feedback Loop Diagrammatic Interception Model, adapted from Sharma et al. (2022)	13
Figure 3 System Architecture Diagram: Three-Layer Edge Computing Design ...	34

Chapter 1

INTRODUCTION

Doomscrolling refers to the compulsive consumption of distressing or negative online content despite its adverse emotional and behavioral effects (Sharma et al., 2022). Mainstream device- and platform-level digital well-being tools still emphasize screen-time summaries, app timers, or break reminders instead of session-level interpretation of fast-changing feed content (Apple Inc., 2025; Google, 2024; Mosseri, 2021; Rahmillah et al., 2023; TikTok, 2024). This thesis addresses that gap through the design, development, and evaluation of the **Heuristic Risk-State Estimation System** — a privacy-preserving Android app that performs on-device risk estimation using behavioral signals, a text-first sentiment path, and an on-device Vision-Language Model (VLM) fallback for no-text items. Here, “detection” refers to computational risk estimation from observable proxies, not diagnostic classification.

The system combines threshold-based heuristics, interpretable fuzzy rule-based inference (Pickering et al., 2025; Vashishtha et al., 2023), a VADER-compatible text-first sentiment pipeline for low-resource code-mixed text (Hutto & Gilbert, 2014; Mohammed & Prasad, 2023; Nazir et al., 2026), and an on-device no-text VLM fallback implemented with Moon-dream 0.5B (Das & Singh, 2023; Sharshar et al., 2025). All processing occurs locally on the user’s device to preserve privacy under real-world mobile constraints. It is evaluated as a bounded software-engineering study, not a clinical trial, and focuses on practical usefulness, software quality, user acceptance, expert review, baseline convergent association, and short-term behavioral differences during a two-week field deployment.

1.1. Background of the Study

In recent years, doomscrolling has emerged as a behavioral concern in algorithm-driven social media environments. Short-form platforms such as TikTok, Facebook Reels, and Instagram Reels are built around rapid, recommendation-driven streams of content, and prior studies associate heavier or more compulsive use with attention-related strain, procrastination, reduced mindfulness, anxiety-related symptoms, and related negative outcomes (Canila et al., 2023; Cardoso et al., 2024; Hawwa et al., 2025; Qin et al., 2022; Taskin et al., 2024; Zhang, 2023).

Existing digital well-being tools address this problem only partly. Prior work shows that the emotional effects of scrolling depend on what users are exposed to, not just how long they scroll (Buchanan et al., 2021). Mainstream responses still center on time limits, usage summaries, or fixed break reminders, while app-level interventions reviewed in the literature remain heterogeneous and do not foreground privacy-by-design as a core requirement (Apple Inc., 2025; Google, 2024; Mosseri, 2021; Rahmillah et al., 2023; Tewari, 2023; TikTok, 2024). This suggests a narrower engineering gap for the present study: whether a private mobile tool can move beyond time alone by combining observable interaction patterns with a text-first sentiment proxy while still handling no-text items through an on-device visual fallback.

Mindfulness-oriented digital interventions provide limited support for brief app-delivered self-regulation prompts (Mitsea et al., 2023). Just-in-time smartphone interventions have also shown that smartphone-based support can be delivered in the moment, with feasibility and acceptability depending on timing, receptivity, and burden management (Mair et al., 2022; Roffarello & De Russis, 2021; Teepe et al., 2021; Yang et al., 2023). These works support in-the-moment prompting, but they do not show whether one system can continuously monitor short-form video behavior, process extracted text locally, resolve no-text

items through the VLM fallback, and issue adaptive prompts without cloud transmission or remote dashboards.

In the Philippine context, local peer-reviewed studies indicate that doomscrolling, problematic social-media use, and digital well-being are relevant concerns, although the strongest available evidence remains student-focused (Canila et al., 2023; Cardoso et al., 2024; Mike Jeizy P et al., 2024). Nearby adult Philippine contexts are more indirect, including social-media use among older adults and anxiety-related social-media exposure findings in adult workers (Castillo et al., 2022; Zamora et al., 2021). The local literature therefore supports contextual relevance, but not adult prevalence, validated thresholds, or expected intervention effects. Against that backdrop, the study examines whether the system can combine behavioral monitoring, a text-first sentiment proxy, the no-text VLM fallback, and adaptive prompting within a privacy-by-design architecture, then be evaluated for usability, reliability, and practical usefulness among adult Filipino users.

1.2. Statement of the Problem

This study addresses the gap left by time-based or platform-controlled digital well-being tools by asking whether a privacy-preserving Android app can be designed, developed, and evaluated for doomscrolling-related risk estimation on short-form video platforms. Current responses remain fragmented: platform-native tools are usually time-based, while related research analogues rarely combine local monitoring, content-sensitive estimation, and adaptive prompting in one deployable adult short-form video system. Because the available local literature does not yet support calibrated estimator thresholds or clinical interpretation, the problem is framed as a software-engineering question. Specifically, it seeks to answer the following questions:

1. What privacy-preserving mobile architecture and estimation framework can support doomscrolling-related risk estimation on the target short-form video

platforms using behavioral indicators and sentiment-related indicators when reliably resolvable, with 2-input behavioral fallback for sentiment-unreliable sessions?

2. What short-term Week 1-to-Week 2 changes are observed in selected logged usage metrics and self-reported doomscrolling scores between the intervention group and the logging-only control group, and within the intervention group across the same period?
3. What baseline convergent association exists between the fixed-prior Week 1 Doomscroll Severity Index (DSI) and participants' self-reported doomscrolling scores among participants with at least three sentiment-reliable Week 1 sessions?
4. How do users evaluate the system using the ISO/IEC 25010 software quality model and the Technology Acceptance Model (TAM)?
5. How do subject matter experts evaluate the system's technical design, privacy safeguards, heuristic logic, and intervention structure?

1.3. Significance of the Study

Contribution to the Body of Knowledge: This study contributes a mobile-systems design and pilot-scale field evaluation of one app that integrates behavioral monitoring, a text-first sentiment proxy, a no-text VLM fallback, adaptive prompting, and on-device privacy. Its contribution is an engineering evaluation framework for feasibility, acceptability, and estimator plausibility. The **Minimum Viable Lexicon (MVL)** workflow supports low-resource, code-mixed text handling within the system rather than functioning as a separate lexicon-development study.

Beyond its academic contribution, the study may offer value to the following stakeholders:

Adult users of short-form video platforms may benefit from a consent-based self-monitoring application that supports reflective interruption of problematic scrolling behavior without sending raw content to external servers.

Software developers and mobile-computing researchers may reuse the app’s architecture, hybrid text/VLM proxy design, and fallback logic when building privacy-preserving digital well-being systems for code-mixed social media environments.

Digital well-being researchers may use the study as an example of how behavioral logging, lightweight sentiment analysis, and adaptive prompting can be combined in a mobile field evaluation.

Educational institutions and digital well-being advocates may use the findings to support responsible discussions about privacy-preserving, user-directed self-regulation tools rather than surveillance-oriented monitoring systems.

1.4. Objectives of the Study

General Objective:

To design, develop, and evaluate the proposed system as a privacy-preserving, on-device Android application for doomscrolling-related risk estimation on short-form video platforms.

Specific Objectives:

1. To design and develop the system to estimate doomscrolling-related risk using behavioral indicators and sentiment-related indicators when reliably resolvable, with 2-input behavioral fallback for sentiment-unreliable sessions.
2. To determine the short-term Week 1-to-Week 2 changes in selected logged usage metrics and self-reported doomscrolling scores between the intervention group and the logging-only control group, and within the intervention group across the same period.

3. To determine the baseline convergent association between the fixed-prior Week 1 Doomscroll Severity Index (DSI) and participants' self-reported doomscrolling scores among participants with at least three sentiment-reliable Week 1 sessions.
4. To evaluate the system from the users' perspective using the ISO/IEC 25010 software quality model and TAM.
5. To obtain subject matter expert evaluation of the system's technical design, privacy safeguards, heuristic logic, and intervention structure.

1.5. Scope and Delimitations

This study covers the design, development, and evaluation of the system. The scope centers on the software artifact, its risk-estimation variables, and the two-week field evaluation.

Scope:

- Development of a native Android application using modern software engineering practices.
- Implementation and evaluation cover one deployable system and a pilot-scale two-week field study.
- Use of the Android Accessibility Service API for text extraction, interaction-event monitoring, and screenshot-assisted no-text routing through `AccessibilityService.takeScreenshot`, with related interaction signals used to estimate session duration, compute video dwell time from content transitions, and process sentiment-related indicators from visible text or no-text visual items. Swipe counts are retained only as supporting logs.
- Design of a modular architecture with separation between data collection, processing, and user interface components.

- Implementation of the core estimation framework using threshold-based rules, text-first VADER sentiment analysis, an on-device no-text VLM fallback using Moondream 0.5B, and fuzzy-logic inference, with unresolved cases handled through sentiment-unreliable classification and 2-input behavioral fallback.
- The system is designed for TikTok, Facebook Reels, and Instagram Reels as target platforms, although actual platform exposure during field deployment may vary by participant and empirical claims are limited to platforms that yield stable extraction.
- On-device processing, where raw text and temporary no-text screen frames are processed locally in RAM and discarded after scoring, while only aggregate local metrics and configuration data are retained on-device.
- A user-directed design in which risk estimates and prompts are delivered only to the consenting user on their own device, with no remote administrator dashboard.
- Evaluation uses a pilot-scale two-week deployment: all participants complete a one-week baseline logging phase, after which the intervention group receives adaptive prompts during Week 2 while the control group continues logging-only.
- Evaluation uses the ISO/IEC 25010 Software Quality Model (Functional Suitability, Performance Efficiency, Usability, Reliability), the Technology Acceptance Model (Perceived Usefulness, Perceived Ease of Use), subject matter expert review, and short-term observed differences under the study conditions.

Delimitations:

- The study targets Filipino Android users aged 18 years and above who actively use at least one of the target short-form video platforms. Its external validity does not extend to minors, iOS users, or users who primarily consume long-form or non-short-form content.

- Field deployment is limited to Android devices that can install the study build, maintain the required permissions, and keep the monitoring service sufficiently stable during the evaluation window.
- The risk-estimation framework addresses only short-form video platforms and does not cover long-form video, news websites, general web browsing, or other social media formats outside TikTok, Facebook Reels, and Instagram Reels.
- The study is not designed as a clinical diagnostic system, medical intervention, psychotherapy tool, or long-term habit-formation trial.
- The study does not claim to prevent doomscrolling. It evaluates a heuristic risk-estimation system and observes short-term behavioral differences between intervention and control groups.
- The study does not attempt multi-site rollout, large-sample validation, or extensive model optimization.
- The study excludes covert monitoring, employer surveillance, parental-control workflows, remote dashboards, and third-party alerting mechanisms.

1.6. Limitations

This study evaluates software behavior and short-term usage differences — not clinical outcomes. Changes in logged metrics should not be interpreted as reductions in anxiety, distress, or mental health conditions.

The following limitations describe constraints beyond the researchers’ control that may affect the study’s findings and generalizability:

- The use of Android Accessibility Service to read on-screen content from third-party applications may create platform-policy concerns. Ethical risk is reduced through explicit user consent, local processing, and the absence of third-party

surveillance workflows, but policy changes remain outside the researchers' control.

- The Accessibility Service approach is fragile by nature. Changes in the user interface of TikTok, Facebook, or Instagram may disrupt extraction logic, and Android device manufacturers may aggressively terminate background services, affecting data continuity.
- Continuous monitoring may affect battery life, thermal behavior, and background stability on some devices, which can in turn influence both system reliability and user acceptance.
- The system does not perform continuous full-resolution video or audio analysis. It relies on behavioral logs, accessibility-extracted text, and sparse transient on-device screen-frame sampling for no-text items. The system therefore still estimates risk from observable proxies rather than from full multimodal understanding of the viewed media.
- Although the deployable app supports Android 8.0 (API 26) and higher, the screenshot-assisted no-text VLM path depends on `AccessibilityService.takeScreenshot`, which is available only on Android 11 (API 30) and higher. On older compatible devices, no-text visual resolution degrades to the sentiment-unreliable / 2-input fallback path when screenshot capture is unavailable.
- Video duration is not consistently exposed across all target applications through the Accessibility Service. For this reason, the system relies on video dwell time and related interaction patterns rather than exact platform-reported video length.
- Because Android Accessibility APIs do not expose a guaranteed content-level identifier for third-party short-form feed items, transition detection remains

heuristic and may still mis-segment some consecutive items, especially when they are text-barren or visually similar.

- The baseline sentiment engine is English-centric. Although the system extends it with a limited Filipino/Taglish lexicon, unsupported dialects, slang drift, and code-mixed text may still reduce sentiment reliability.
- Sessions with high out-of-vocabulary text, or no-text items that still cannot be resolved reliably after VLM routing, are treated as sentiment-unreliable and excluded from some sentiment-dependent analyses, which may reduce the effective analyzable sample for some outcomes.
- Self-reported profile data and Doomscrolling Scale responses remain subject to recall bias and social desirability bias.
- Because recruitment is limited to consenting participants willing to install a monitoring application, the sample may overrepresent users who are already receptive to self-regulation tools.
- Although a logging-only control group is included to isolate prompt effects from monitoring awareness and natural fluctuation, the study remains non-blinded: intervention-group participants know they are receiving prompts, which may introduce expectancy or novelty effects beyond the prompt mechanism itself.
- For the intervention group, prompt-excluded active-use metrics can be mechanically lower by design when prompts occur because prompt-display time is removed from session duration and dwell-time calculations. For this reason, Chapter 3 treats raw elapsed session duration and raw elapsed dwell time as the primary behavioral comparison, while prompt-excluded metrics are kept only as supplementary traces of prompt-interrupted use.

- The two-week deployment window captures only short-term behavior. It is insufficient for demonstrating long-term habit formation, retention, or sustained behavioral change.
- Participants know they are being monitored, which introduces Hawthorne-effect risk and may reduce the naturalism of observed behavior.

1.7. Conceptual Framework

The study uses the **Input-Process-Output (IPO) model** to illustrate the implemented runtime architecture and data flow for heuristic risk estimation (Laudon & Laudon, 2022).




Figure 1: Conceptual Framework of the Study using IPO Model

The IPO model here refers to the runtime application pipeline rather than the full research workflow. During Week 2, a **feedback loop** operates once the live duration gate is met and the live RiskScore enters the Warning or Critical bands: the app selects an intervention level (L1 awareness, L2 pause, or L3 breathing exercise), closes the current interval, and begins a new one only if the user returns to the target platform. When sufficient reliable baseline sessions exist, the live prompt engine may personalize Session Duration and NSD memberships from Week 1 quantiles while Video Dwell Time remains fixed. Usable caption or visible-comment text stays on the text path; no-text items are routed

through `AccessibilityService.takeScreenshot` and the VLM fallback; unresolved cases degrade to 2-input behavioral inference with conservative L1/L2 prompts.

1.8. Theoretical Framework

The system is informed by the **Doomscrolling Feedback Loop Model** derived from Sharma et al. (2022). The model is used as a design guide for content exposure, prolonged engagement, and interruption points, but the study does not directly measure the internal emotional states described by the model.

The model conceptualizes doomscrolling as a three-phase cycle:

1. **Antecedents or triggers:** Users are drawn toward distressing or uncertainty-inducing content by factors such as anxiety, fear of missing out, or crisis-related information needs.
2. **Behavior:** Users engage in persistent, repetitive scrolling through negative or emotionally charged content.
3. **Outcome:** Exposure to such content contributes to distress or negative affect, which may in turn reinforce further scrolling.

This framework informs the system’s variable selection. The system estimates risk at the behavior and exposure level by observing video dwell time, session duration, and NSD as proxies for heightened maladaptive engagement. Swipe or scroll events are used only as transition signals to determine when one short-form item ends and another begins, allowing video dwell-time computation rather than serving as a separate estimator input. Its adaptive prompts are intended to interrupt potentially escalating patterns before exposure becomes more sustained.

The framework thus justifies using observable proxies and non-clinical prompts. Evaluation is guided by ISO/IEC 25010 and TAM, while the baseline-to-intervention design supports short-term behavioral comparison.




Figure 2: Doomscrolling Feedback Loop Diagrammatic Interception Model, adapted from Sharma et al. (2022)

The system is therefore a **heuristic risk-estimation tool** that uses behavioral and exposure proxies rather than direct measures of internal psychological states.

1.9. Definition of Terms

Adaptive Digital Mindfulness Prompts - Non-clinical system-generated interventions, such as awareness notifications, pause prompts, or short breathing exercises, that are triggered according to estimated risk severity to encourage reflective interruption of scrolling behavior.

Digital Mindfulness - A practice of intentional and self-aware technology use that promotes attention, presence, and self-regulation in digital environments (Aggarwal et al., 2024).

Doomscrolling - The compulsive and continuous consumption of negative or distressing content on digital platforms, particularly in ways that may reinforce anxiety, vigilance, or maladaptive engagement (Sharma et al., 2022).

Doomscrolling-Related Risk Estimation - The computational estimation of elevated scrolling risk using observable behavioral and sentiment-related proxies rather than direct measurement of internal psychological states.

Edge Computing - A computing approach in which data processing occurs on or near the data source - in this study, on the user's mobile device - rather than in remote cloud infrastructure, thereby reducing latency and privacy risks (Swathi et al., 2025).

Fuzzy Logic System - A computational approach that uses partial membership and rule-based reasoning to model ambiguous behavioral patterns, allowing graded risk estimation instead of binary classification through interpretable linguistic rules (Pickering et al., 2025; Vashishtha et al., 2023).

Heuristic Risk-State Estimation System - The Android application developed in this study for heuristic doomscrolling-related risk estimation and adaptive prompting.

ISO/IEC 25010 - An international software quality model used in this study to evaluate selected quality characteristics of the system, specifically Functional Suitability, Performance Efficiency, Usability, and Reliability (International Organization for Standardization, 2023).

Negative Sentiment Density (NSD) - The proportion of resolvable content units within a session that are classified as negative. In this study, it is computed from analyzable caption and visible-comment text units that pass the text-side reliability screen together with no-text items resolved through the VLM fallback. It is used as a limited proxy for negative exposure.

Technology Acceptance Model (TAM) - A framework used to assess user acceptance of digital systems, particularly through the constructs of Perceived Usefulness and Perceived

Ease of Use, and still commonly applied in mHealth acceptability evaluation (Adnan et al., 2025).

VADER Sentiment Analysis - Valence Aware Dictionary and sEntiment Reasoner, a lexicon- and rule-based sentiment analysis method for social media text that is adapted in this study for limited code-mixed use (Hutto & Gilbert, 2014).

Vision-Language Model (VLM) Fallback - A limited on-device visual-analysis path used in this study when a viewed short-form item has no usable caption or visible comments. It uses the Moondream 0.5B model together with `AccessibilityService.takeScreenshot` and transient RAM-only frame processing to estimate whether a no-text item contributes negative exposure.

Video Dwell Time - The time interval spent on a viewed short-form video before the user swipes away or otherwise transitions, computed from heuristically inferred content-transition signals derived from accessibility events and used in this study as one behavioral proxy for scrolling engagement.

Very Large Online Platforms (VLOPs) - Large digital platforms, such as Facebook, Instagram, and TikTok, that operate at massive user scale and strongly shape online content exposure and engagement patterns (Chen et al., 2024).

Session Duration - The continuous period of active engagement with a target short-form video platform, as determined by interaction signals and the session-management rules described in Chapter 3.

Sentiment-Unreliable Session - A session in which item-level negativity cannot be derived reliably for the main analysis after the available text path or no-text VLM path has been attempted, such as sessions dominated by high-OOV text or by unresolved no-text items. Such sessions are excluded from the primary NSD- and DSI-based analyses.

Swipe/Scroll Transition Signals - Accessibility-detected navigation events used to heuristically infer when one short-form video ends and another begins, with text verification

applied when usable caption text exists, enabling the computation of video dwell time and the logging of supporting interaction counts.

Chapter 2

REVIEW OF RELATED LITERATURE AND STUDIES

This chapter reviews peer-reviewed literature and empirical studies relevant to doomscrolling, short-form platform engagement, digital mindfulness, privacy-preserving mobile monitoring, sentiment analysis, and software evaluation. Journalistic, policy, platform-issued, and product-documentation materials are used only for contextual background or feature verification. The review is organized to clarify the study’s evidence base, technical grounding, and remaining gaps.

2.1. Foreign Literatures

This section presents selected foreign literature relevant to doomscrolling, short-form platform design, digital mindfulness, and privacy-preserving mobile systems.

2.1.1. Doomscrolling, Platform Design, and Compulsive Engagement

Recent literature treats doomscrolling in short-form feeds as a joint product of user vulnerability and platform design rather than as a simple matter of screen-time volume. Rodrigues (2022) frames doomscrolling as compulsive information seeking that can intensify anxiety and helplessness, while Zenone et al. (2021) show that recommendation-driven environments can amplify harmful exposure and addictive-use patterns. Zhang (2023) complements that view by showing that TikTok is woven into users’ routines as entertainment, information, and social reference, which helps explain why these feeds can sustain engagement so effectively.

These works suggest that time-only intervention logic is too coarse for the present problem. A defensible intervention system should monitor sustained interaction patterns within algorithmic short-form feeds rather than rely only on coarse daily time limits.

2.1.2. Digital Mindfulness and Technology-Assisted Self-Regulation

Recent digital well-being reviews converge on a narrow design implication: technology-assisted self-regulation support is more acceptable when it is brief, timely, and low-burden. Mitsea et al. (2023) provide limited support for digitally assisted mindfulness as one self-regulation strategy, while Antezana et al. (2022) show that engagement improves when interventions remain lightweight and relevant to the user’s immediate context. Together, they support short adaptive prompts as low-burden, non-clinical interruption mechanisms.

2.1.3. Privacy, Persuasive Design, and Ethical Considerations in Digital Well-Being

Privacy-oriented digital well-being literature also constrains what a defensible system can do. Tewari (2023) argues that privacy protections should be designed into mobile health architectures from the start rather than added after collection logic is already fixed. That point is central to the present design: behavioral monitoring, sentiment processing, and prompting should occur without exporting raw user content to external services. This literature therefore supports local processing and aggregate-only retention.

2.1.4. Edge Computing and On-Device Vision-Language Fallbacks

Recent literature on privacy-preserving mobile analysis converges on a clear constraint: if visual inference is used at all, it should remain inside the device’s privacy boundary. Sending screenshots to cloud-based multimodal APIs would expose sensitive user data to external servers. Sharshar et al. (2025) support the feasibility of edge-deployed Vision-Language Models under network and latency constraints, while Lee et al. (2024)

show why compact transformer-based vision models remain a practical design priority on mobile and edge devices. Broader multimodal sentiment and affective-computing reviews likewise show that visual channels can add information when text is absent, although interpretation remains limited (Cortiñas-Lorenzo & Lacey, 2024; Das & Singh, 2023; Johnson et al., 2025). This literature supports the app’s on-device VLM fallback for no-text items, with Moondream 0.5B used as a small deployment-fit model.

2.2. Local Literatures

This section presents selected local literature and contextual materials relevant to doomscrolling, excessive social media use, and digital well-being in the Philippine setting. Because adult Philippine doomscrolling evidence remains limited, peer-reviewed local studies are used as nearby evidence of attention, scrolling, and self-regulation concerns, while journalistic or policy-oriented sources serve only as contextual background.

2.2.1. Local Discourse and Platform Concern

Philippine background materials show that doomscrolling and digital well-being had already entered local public discourse by 2021 Lanuza et al. (2021); Panaligan (2021). They establish contextual relevance, not adult behavioral prevalence, intervention effectiveness, or estimator validity.

2.2.2. Limits of the Local Evidence Base

The local empirical base remains uneven and still depends mainly on student-centered studies discussed later in this chapter. Recent Philippine work outside student-only samples is useful but indirect: Cleofas et al. (2022), Zamora et al. (2021), and Castillo et al. (2022) broaden the adult context by showing that social-media use intersects with mental well-being, anxiety-related susceptibility, or social connectedness across different Filipino

populations. Even so, the local literature does not yet provide adult doomscrolling field evidence, validated cutoffs, or intervention-effect benchmarks for a privacy-preserving mobile system.

2.3. Interim Synthesis of Literature

The reviewed foreign and local literature establishes three points. First, doomscrolling is a recognizable behavioral construct linked to distress-related and attention-related concerns, and platform design appears to help sustain engagement (Hawwa et al., 2025; Qin et al., 2022; Taskin et al., 2024). Second, existing digital well-being responses remain fragmented: mainstream device- and platform-level tools still emphasize summaries, timers, or break reminders rather than user-controlled, locally processed, content-sensitive estimation (Apple Inc., 2025; Google, 2024; Mosseri, 2021; Rahmillah et al., 2023; Tewari, 2023; TikTok, 2024). Third, Philippine literature supports local relevance and provides some adult context, but not enough adult field evidence to justify prevalence claims, threshold values, or expected intervention effects.

Together, these sources justify an engineering problem: combining behavioral signals, visible text, and a no-text VLM fallback in one privacy-preserving mobile system. They do not support text sentiment as a complete proxy for multimodal short-form feeds or provide strong local adult evidence for universal estimator parameters. Accordingly, the study uses a hybrid exposure signal in which analyzable text is scored through the text path and no-text items are routed through the VLM fallback.

2.4. Foreign Studies

This section presents a review of related studies that provide empirical evidence on doomscrolling, compulsive social media use, and their psychological and behavioral impacts.

2.4.1. Psychometric Validation of Doomscrolling as a Measurable Construct

Sharma et al. (2022) establish doomscrolling as a measurable self-report construct, while Satıcı et al. (2023) extend that psychometric work and support a shorter 4-item form. Together, these studies justify using self-reported doomscrolling as a comparison anchor, not as session-level ground truth. Chapter 3 therefore treats the scale as a baseline convergent-association reference rather than as a direct label for individual scrolling sessions.

2.4.2. Psychological Mechanisms and Consequences of Doomscrolling

Taskin et al. (2024) and Hawwa et al. (2025) both position doomscrolling as part of a wider pathway linking platform engagement to poorer well-being outcomes. Taskin et al. (2024) highlight reduced mindfulness and secondary traumatic stress as mediating mechanisms, while Hawwa et al. (2025) connect doomscrolling to the relationship between social media addiction and anxiety. Their shared implication is that a pure time-only measure is insufficient, yet neither study validates a single behavioral proxy or mobile risk-estimation rule by itself.

2.4.3. Platform Addiction and Algorithmic Engagement

Qin et al. (2022) show that TikTok’s platform characteristics can intensify flow experience and addictive use. This supports the thesis’s focus on short-form mechanics and sustained engagement rather than on user weakness alone.

2.4.4. Computational Approaches to Detecting Harmful Digital Behavior

Zannettou et al. (2024) show that engagement with TikTok recommendation streams can be studied from donated behavioral traces, while Lokeshkumar et al. (2021) show that text and activity patterns can be processed computationally as risk-relevant indicators. These studies serve as methodological analogies: they support technical plausibility for

trace-based analysis, but not direct precedent for doomscrolling estimation on short-form platforms or validation of the specific proxies and privacy constraints used here.

2.4.5. Fuzzy Logic as an Inference Engine for Mobile Behavioral Classification

Recent reviews show that fuzzy rule-based systems are useful when the target construct is gradual rather than binary and when the inference process must remain interpretable to human reviewers (Pickering et al., 2025; Vashishtha et al., 2023). Recent work on interpretable fuzzy-model design also continues to treat simple triangular or shoulder-style membership sets as practical, auditable choices when labeled session-level training data are not available (Khairuddin et al., 2021; Porebski, 2022). Recent fuzzy-design work also describes symmetric grid partitions and strong fuzzy partitions as interpretable starting structures when domain-specific numeric cutoffs are unavailable (Azam et al., 2021; Casalino et al., 2022; Khairuddin et al., 2021). Together, this literature supports using fuzzy logic and an interpretable starting structure for Negative Sentiment Density, where no universal percentage threshold set is established.

2.4.6. Data Extraction and Privacy-Preserving Edge Computing

Lee et al. (2022) show that Android Accessibility Service can expose fine-grained interaction data unavailable through standard APIs, but they also make clear that such access is privacy-sensitive. Swathi et al. (2025) complement that point by supporting edge-centric monitoring as a privacy-preserving and responsive architecture. Together, these studies justify the architectural stance that rich interaction data, if collected at all, should be processed locally and reduced to aggregate outputs rather than exported as raw content.

2.4.7. Social Media Use, Sustained Attention, and Digital Well-Being Interventions

Rajeswari & Vakkil (2024) and Roffarello & De Russis (2021) support different parts of the present design. Rajeswari & Vakkil (2024) link heavier social media engagement with poorer sustained attention, which supports the general relevance of prolonged use and viewing cadence as behavioral signals. Roffarello & De Russis (2021) show that smartphone habits can be monitored in real time and mitigated through proactive just-in-time reminders, which supports the feasibility of lightweight adaptive intervention on the device. Together, they justify attention to both continuous engagement patterns and in-the-moment interruption, though not direct validation of the present doomscrolling estimator. Recent measurement studies also help bound the fuzzy time-based inputs, even though they do not validate doomscrolling-specific cutoffs. Muise et al. (2024) show that smartphone content exposure often occurs in bursts of only a few seconds, Cho et al. (2021) treat sessions shorter than 5 seconds as likely mistaken openings and use 30 seconds as a short-duration divider in feature-level social-media logging, and Tian et al. (2021) report that 93.7% of mobile app usage in their large-scale logs lasts less than 10 minutes. Recent smartphone-logging work that reconstructs sessions from app-event streams also continues to merge adjacent app uses into one session when inter-app gaps remain within about 45 seconds (Ahmed & Ahmed, 2023; Chen et al., 2023). These studies justify distinguishing immediate skips, brief but nontrivial viewing, clearly sustained use, and very short app-switch interruptions instead of treating all exposure as one undifferentiated time scale. Intervention-timing studies point in the same direction. Terzimehić & Aragon-Hahner (2022) report that sessions shorter than 10 minutes are less prone to regret, with alternative offline activities becoming more salient in the 10-20-minute and >20-minute ranges. Rixen et al. (2023) treat 10+ minute infinite-scrolling episodes as long sessions, Meinhardt et al. (2025) deploy an infinite-scroll intervention after 15 minutes of continuous scrolling, Ismail & Al Thani (2022) use personalized prompts after 40 minutes of continuous sitting, and

Ikegaya et al. (2025) show that week-1-derived personalized intervention criteria outperform uniform criteria in the first hour after prompting. Recent JITAI reviews add the same design implication at the framework level: decision points and rules should be empirically grounded, burden-aware, and personalized when possible rather than fixed by one universal cutoff (Fiedler et al., 2024; Genugten et al., 2025; Hsu et al., 2025). Together, these studies support treating sub-10-minute use as a lower-risk baseline region, recognizing 15 minutes as a defensible live-prompt entry point for continuous infinite scrolling, and deriving Week 2 prompt criteria from each participant’s Week 1 baseline.

2.4.8. Technology Acceptance Frameworks

Recent mHealth acceptability reviews continue to treat TAM as a practical evaluation framework, with Perceived Usefulness and Perceived Ease of Use retained as core acceptance constructs (Adnan et al., 2025). That supports the inclusion of TAM alongside ISO/IEC 25010 in the post-usage evaluation.

2.4.9. Computational Constraints in Code-Mixed Sentiment Analysis

Analyzing Taglish on mobile devices presents two challenges: computational load and high Out-of-Vocabulary (OOV) rates. Recent low-resource and code-mixed sentiment work shows that transformer approaches can be data-hungry or resource-intensive, while lexicon-based and hybrid methods remain relevant when reproducibility, low overhead, and resource scarcity matter (Hussain et al., 2025; Mohammed & Prasad, 2023; Nazir et al., 2026; Tho et al., 2021). Work on lexicon curation for low-resource social-media text also shows that slang, emoticons, and local usage patterns often require targeted valence tuning rather than direct reuse of generic sentiment resources (Hashmi et al., 2024; Khan et al., 2025; Perera & Caldera, 2024; Wijayanti & Arisal, 2021). For that reason, the study adopts

a limited VADER + **Minimum Viable Lexicon** (MVL) strategy for code-mixed sentiment analysis on mobile devices.

2.4.10. Addressing the Affective Gap Through an On-Device VLM Fallback

When a short-form item contains no usable caption or visible comments, text-only sentiment cannot characterize that item at all. Reviews of multimodal sentiment analysis show why visual channels become relevant once text is absent (Das & Singh, 2023), while user-generated-video emotion studies show that keyframe-based reduction can support practical visual inference from video without requiring dense frame-by-frame processing (Wei et al., 2021; Zhang & Xu, 2023). Additional recent work shows that sparse visual sampling can remain useful under resource or privacy constraints, including sparse sampling with majority voting on short utterance clips (Sharma et al., 2023), privacy-compliant group emotion recognition using only 5 uniformly distributed frames on 5-second videos (Augusma et al., 2023), and sparse sampling regimes in recent VLM long-video evaluation under context-length constraints (Qu et al., 2025). These studies justify routing no-text items through the VLM fallback, using Moondream 0.5B as a small deployment-fit model with an on-demand screenshot capture rule rather than continuous video analysis. They do not validate the exact prompt design or label set, and they do not support full multimodal affective understanding.

2.5. Local Studies

2.5.1. Doomscrolling and Student Outcomes in the Philippines

Mike Jeizy P et al. (2024) provide the closest local study to the present topic by describing doomscrolling-related experiences among Filipino students. Bautista et al. (2024), Canila et al. (2023), Ababat et al. (2024), and Cleofas et al. (2022) reinforce the

same local signal: heavier or more problematic social media use is linked to boredom proneness, distraction, procrastination, attention-related strain, or poorer mental well-being among Filipino students and young adults. These studies establish Philippine relevance but remain student-centered, self-report-heavy, or qualitative, so they support cautious adult field evaluation rather than adult norms or threshold values.

2.5.2. Social Media, Reading Competence, and Attention Span

Gagalang (2021) and Cardoso et al. (2024) together support the local relevance of attention-related harms associated with heavy digital-platform use, especially among students. They help explain why viewing cadence and sustained use are reasonable behavioral dimensions to observe, even though they do not validate the estimator thresholds used for an adult sample.

2.5.3. Cognitive Benefits of Digital Detox Interventions

Lim et al. (2025) show that a two-week digital detox can improve attention and memory among Filipino Grade 12 students. The study is relevant mainly because it makes a two-week observation window plausible, not because it is a direct precedent for the present prompt-based intervention.

2.5.4. Local Computational Work on Filipino Sentiment Analysis

Co et al. (2022) and Cruz & Balahadia (2022) show that Filipino or English-Filipino public text can be analyzed with lexicon-based or lightweight sentiment pipelines, but they also make clear that domain adaptation remains necessary when language use shifts across settings. These studies support the narrower claim that Filipino-facing sentiment processing is technically workable in the local setting, provided the lexicon is treated as task-specific and incomplete rather than as a universal sentiment resource.

2.5.5. Software Quality Evaluation in Philippine Educational Technology

For software evaluation, the stronger methodological anchor is the official ISO/IEC 25010 product quality model itself rather than a single local implementation paper (International Organization for Standardization, 2023). The standard supports the choice of selected quality characteristics, while the survey wording used in Chapter 3 remains researcher-developed and subject to expert review.

2.6. Integrated Synthesis and Research Gap

2.6.1. Comparative Analysis of Existing Digital Well-Being Tools

To contextualize the research gap identified in this review, the following table summarizes the capabilities of existing digital well-being tools and platform-level features. The comparison is based on five key requirements derived from the preceding literature. Feature characterizations for Apple Screen Time, Google Digital Wellbeing, TikTok, and Instagram are drawn from their official product or support pages (Apple Inc., 2025; Google, 2024; Mosseri, 2021; TikTok, 2024). The broader third-party app row is grounded in the multimethod review by Rahmillah et al. (2023).

Operational Definitions for Comparison:

- **Behavioral Monitoring:** [Real-time] = continuous background monitoring of specific actions; [Aggregate] = summary views or timer-based monitoring without item-level sensing; [Platform-only] = monitoring limited to one platform's own activity tools.
- **Text-Visible Sentiment Proxy:** [✓] = analyzes the emotional tone of extracted text content; [×] = treats all viewed content equally regardless of text tone.
- **Adaptive Intervention:** [✓] = intervention type changes based on estimated severity; [Partial] = fixed caps, timers, or break nudges that do not estimate

content or risk; [×] = monitoring only or fixed reminders without severity adaptation.

- **On-Device Processing:** [✓] = core monitoring or intervention logic is user-controlled and executed locally on the device; [N/A] = feature is platform-managed and not exposed as a user-controlled local processing pipeline.
- **Platform Scope:** [System-wide] = works across many apps on the device; [Single-platform] = limited to one service; [Varies] = depends on the specific third-party app; [Specific targets] = limited to named study platforms.

Table 1. Comparative Analysis of Digital Well-Being Tools and the Proposed System

Tool / System	Behavioral Monitoring	Text-Visible Sentiment Proxy	Adaptive Intervention	On-Device Processing	Platform Scope
Apple Screen Time	Aggregate	×	Partial	✓	System-wide
Google Digital Wellbeing	Aggregate	×	Partial	✓	System-wide
TikTok screen-time tools	Platform-only	×	Partial	N/A	Single-platform: TikTok
Instagram daily limit / break reminders	Platform-only	×	×	N/A	Single-platform: Instagram
Rahmillah et al. (2023) apps	Varies	×	Partial	Varies	Varies
Proposed System	Real-time	✓	✓	✓	Specific targets: TikTok, Facebook Reels, Instagram Reels

As shown in the table, the reviewed literature and official product pages identify useful summaries, timers, and break mechanisms, but not a single adult short-form video system that combines real-time behavioral monitoring, a text-first sentiment proxy, an implemented no-text VLM fallback, adaptive prompting, and strict local processing. The gap therefore lies in integration and deployment scope: the ingredients appear separately in prior work, but not as one privacy-preserving field-tested artifact under the conditions targeted here. Buchanan et al. (2021) help motivate this search by showing that scrolling outcomes depend on exposure quality rather than time alone, while Roffarello & De Russis (2021) show

that just-in-time mobile interventions are feasible even without the specific doomscrolling construct or item-level negativity proxy used here.

The reviewed studies converge on three themes. First, doomscrolling can be measured through self-report and discussed as a distinct construct, which supports a baseline convergent association check (Satici et al., 2023; Sharma et al., 2022). Second, technical and intervention precedents exist for mixed mobile interventions, interpretable fuzzy inference, privacy-preserving on-device processing, Filipino-facing sentiment analysis, and narrow no-text visual inference on edge devices, although the closest studies remain analogies rather than direct implementations of the app evaluated here (Co et al., 2022; Cruz & Balahadia, 2022; Das & Singh, 2023; Lee et al., 2022; Pickering et al., 2025; Roffarello & De Russis, 2021; Sharshar et al., 2025; Vashishtha et al., 2023; Wei et al., 2021). Third, the local Philippine literature indicates that attention, self-regulation, and problematic-use concerns are locally relevant, that some adult social-media contexts have been studied, that short-term digital well-being observation windows are workable, and that Filipino-facing sentiment processing is technically feasible in local text settings (Castillo et al., 2022; Cleofas et al., 2022; Co et al., 2022; Cruz & Balahadia, 2022; Lim et al., 2025; Mike Jeizy P et al., 2024; Zamora et al., 2021). It does not yet establish adult prevalence, intervention effects, or calibrated threshold values for the present estimator.

A clear **research gap** therefore remains. No reviewed study integrates behavioral logging, a text-first sentiment proxy, an implemented no-text VLM fallback, adaptive prompts, and strict local processing in one privacy-preserving adult short-form video system while also reporting a field evaluation aligned with this thesis. The local adult evidence base is also too limited to justify calibrated estimator thresholds, strong expected-effect claims, or clinical interpretation. Hybrid text-plus-visual sentiment routing is therefore treated as partial coverage of negative exposure, not as a full representation of multimodal short-form video

content. In this study, NSD is defined as a session-level negative-exposure metric computed from resolvable text units and VLM-resolved no-text items. The study addresses these gaps through a two-week pilot field evaluation focused on software behavior, short-term behavioral comparison, user acceptance, expert review, baseline convergent association, and sentiment-reliable coverage reporting.

Chapter 3

METHODOLOGY

This chapter presents the methodology used to design, develop, and evaluate the system. It covers the research design, architecture, development procedure, evaluation method, data collection, respondents, and statistical treatment. Here, the term “detection” refers to heuristic risk-state estimation from observable proxies rather than diagnostic classification.

3.1. Type of Research

This study adopts a **design-and-development research** approach with a **two-group baseline-intervention pilot field evaluation** (Creswell & Creswell, 2022). The design-and-development component focuses on building the software artifact and its heuristic risk-estimation logic, while the field-evaluation component observes the system under real-world conditions.

The study follows a primarily quantitative orientation, with open-ended feedback used as supplementary qualitative context. Automated system logs and structured survey responses are used to assess software quality, user acceptance, expert review, baseline convergent association, and short-term behavioral differences. The study is a software evaluation rather than a clinical trial, and findings are interpreted within the study conditions.

3.2. Research Design

The study uses a **two-group baseline-intervention field design** treated as a **pilot-scale randomized field evaluation**. After consent and study-compatibility screening,

participants are assigned 1:1 to either an **intervention group** or a **logging-only control group** using a concealed permuted-block allocation list with randomly ordered block sizes of 2 and 4. The allocation list is generated before enrollment by a researcher who is not responsible for participant screening and is released only after eligibility is confirmed, which helps preserve balance without making the next assignment easily predictable. All participants then complete an identical one-week baseline logging phase with prompts disabled. During Week 2, the intervention group receives adaptive prompts while the control group continues logging-only.

The study questions are answered through related but distinct methodological parts. Research Question 1 is addressed by the architecture, routing, variable-justification, and algorithm sections. Research Question 2 is addressed by Week 1-to-Week 2 change analyses of selected logged metrics and self-reported doomscrolling scores, comparing the intervention and control groups and using a paired within-intervention comparison as secondary evidence. As a pilot-scale study, these analyses estimate short-term observed differences under the study conditions rather than definitive intervention efficacy. Research Question 3 is addressed by the baseline convergent association between Week 1 DSI and the Doomscrolling Scale. Research Questions 4 and 5 are addressed by the ISO/IEC 25010, TAM, and SME evaluations.

3.3. System Architecture

The system uses a **modular three-layer architecture** to support maintainability, interpretability, and separation of concerns. Its design follows **Privacy by Design**, keeping processing local to the device through an edge-computing approach. The application is organized into the **Data Access Layer**, **Business Logic Layer**, and **Presentation Layer**, allowing data collection, estimation logic, and user-facing components to be developed

with minimal coupling. The architecture also avoids device rooting by relying on Android Accessibility Service for text extraction, interaction-event monitoring, and screenshot-assisted VLM fallback when no usable caption or visible comments exist. The figure below shows this hybrid text/VLM pipeline.




Figure 3: System Architecture Diagram: Three-Layer Edge Computing Design

3.3.1. Data Extraction Strategy

The primary mechanism for data acquisition is the Android **Accessibility Service API**. It lets the system read captions, comments, and interaction events from short-form video apps without requiring device rooting. The same layer also captures the interaction signals needed for session duration, content-transition detection, video dwell-time computation, and supporting swipe logs. When a viewed item has no usable caption or visible comments, the enabled accessibility service uses `AccessibilityService.takeScreenshot` to capture an on-demand transient screen frame for the no-text VLM fallback.

3.3.2. Accessibility Service Implementation

The implementation entails creating a background service that extends the `AccessibilityService` class. Functionally, this allows the system to observe the UI hierarchy of other apps in real-time, effectively perceiving the screen content as a screen reader would for accessibility purposes.

The technical implementation involves the following core components:

- **Service Extension:** Development of a custom service inheriting from `AccessibilityService`.
- **Event Listening:** The service subscribes to specific system events, primarily `AccessibilityEvent.TYPE_WINDOW_CONTENT_CHANGED` and `TYPE_VIEW_SCROLLED`, to detect dynamic content updates.
- **Node Traversal:** Upon capturing an event, the system traverses the hierarchy of `AccessibilityNodeInfo` objects, filtering specifically for `TextView` elements to extract creator captions and any visible user comments already present on-screen as affective reaction signals. The system does not actively open hidden comment panels or trigger additional UI states during extraction.

- **Scroll-Event Disambiguation:** Because `TYPE_VIEW_SCROLLED` events are generated by both feed-level video transitions and within-video comment-panel scrolling, the system does not treat every scroll event as a new item by default. When usable caption text is present, the system applies a content-change verification step before registering a content transition: after each feed-level scroll event, the caption text is re-extracted and compared to the previously recorded caption using exact string equality after whitespace normalization, and a transition is registered only if the caption has changed. When a viewed item has no usable caption or visible comments, however, the system cannot rely on caption comparison; in that no-text case, a feed-level scroll event is treated as a heuristic new-item boundary unless suppressed by comment-panel state or similar non-feed interaction context. As a secondary check, the system tracks comment-panel open/close state via `TYPE_WINDOW_STATE_CHANGED` to suppress transition processing during comment browsing. Comment-panel scrolling, caption expansion, and UI-only animations therefore do not reset the dwell-time interval or increment the swipe count. Comments visible during panel scrolling are still extracted for sentiment analysis and associated with the current video item.

3.3.3. Constraints and Challenges

A significant challenge in this architecture is the requirement to reverse-engineer the UI structure of targeted applications. The extraction logic must identify specific resource IDs (e.g., `com.zhiliaapp.musically:id/caption_text`) to distinguish meaningful content from UI clutter such as buttons or timestamps. This approach has inherent fragility: changes to the target application’s UI structure or IDs by their respective developers may disrupt the scraper’s functionality, representing a known limitation of this methodology.

3.3.4. Edge Computing and Privacy Architecture

The system follows edge computing principles to mitigate privacy risks:

1. **No Persistent Storage of Raw Content:** In the study version, scraped text content and temporary no-text screen frames are processed entirely in RAM and immediately discarded after scoring. Only derived numeric scores, aggregate usage metrics, and local configuration data are retained on-device.
2. **On-Device Processing:** All Fuzzy Logic inference, VADER analysis, and no-text VLM inference occur locally on the user’s device. No content is transmitted to external servers.
3. **User-Initiated Consent:** The Accessibility Service requires explicit user activation through Android system settings. The same enabled service provides the no-text screenshot capability through `android:canTakeScreenshot="true"` and `AccessibilityService.takeScreenshot`, so the study’s operational data-access consent is centered on the Accessibility Service setup.
4. **Aggregate-Only History:** The usage dashboard displays only aggregate statistics (e.g., “negative content exposure: 45%”) rather than storing individual text samples or captured frames.

3.3.5. Extraction Failure Handling

This architecture relies on hardcoded UI resource IDs (e.g., `com.zhiliaapp.musically:id/caption_text`) within the Accessibility Service. Mid-experiment UI updates can restructure the target application’s view hierarchy and disrupt extraction. Such breakage is logged as a **reliability event**, and sessions or participant-periods that cannot yield valid extraction output are excluded from the analyses that require those measures.

3.3.6. Mindfulness Intervention System

Based on digital well-being and smartphone-habit intervention literature, including recent just-in-time adaptive intervention work on timing, feasibility, and burden management (Ismail & Al Thani, 2022; Mair et al., 2022; Roffarello & De Russis, 2021; Teepe et al., 2021; Wang et al., 2023; Yang et al., 2023), the system uses a multi-level adaptive intervention approach:

Level 1 - Awareness Notification: A non-intrusive toast message appears when the RiskScore enters the lower Warning band, informing the user of their session duration.

Level 2 - Pause Prompt: A modal dialog offers the user options to continue, take a break, or view their usage statistics when the RiskScore enters the upper Warning band.

Level 3 - Breathing Exercise: A full-screen overlay presents a brief guided breathing animation (study default: 60 seconds) when the RiskScore enters the Critical band. The breathing overlay is treated as a study-defined pause-and-reset prompt rather than as a therapeutic procedure (Mitsea et al., 2023; Yang et al., 2023). It is intended as behavioral interruption only.

For analytic transparency, time spent inside system-generated modals or overlays is excluded from prompt-excluded active-use metrics. When an intervention appears, the ongoing content-view interval is closed before the prompt begins, and a new interval starts only if the user returns to the target platform afterward. Raw elapsed session time and raw elapsed dwell time are retained separately for the primary behavioral comparison in the evaluation phase.

Usage Dashboard: Shows aggregated usage and sentiment trends over time to support self-reflection and behavioral awareness. It is treated as a secondary reflective feature rather than as the primary intervention mechanism.

3.3.7. Design Justification for Risk-Estimation Variables

This section addresses Research Question 1 by mapping each risk-estimation variable to the Doomscrolling Feedback Loop Model (Sharma et al., 2022) and to the reviewed literature. Variables were retained only if they were conceptually plausible, literature-supported, and observable through Android without rooting.

Table 2. Design Justification Matrix: Theoretical Constructs to System Variables

Theoretical Construct	Observable Proxy	System Variable	Data Source	Unit
Behavior: Persistent, obsessive consumption	Prolonged or rapid navigation through content without meaningful pauses	Video Dwell Time	AccessibilityService (active time between TYPE_VIEW_SCROLLED events, gated by micro-interaction events)	seconds
Behavior: Repetitive rapid feed navigation	Frequency of moving through items during a session	Swipe Count (supporting interaction log)	AccessibilityService view-scroll event stream	count
Behavior: Prolonged maladaptive engagement	Extended time on target apps without voluntary breaks	Session Duration	Accessibility-derived session tracker and app-state events	minutes
Antecedents/ Triggers: Negative content exposure	Proportion of consumed content with negative emotional tone	Negative Sentiment Density (NSD)	Negative-exposure density from analyzable text units and VLM-resolved no-text items	% (0–100)

Video Dwell Time was selected because TikTok and Reels use a discrete, paginated swipe interface rather than continuous scrolling. Dwell time - measured as the duration in seconds between consecutive content-verified transition events — captures how quickly or slowly users move through content. Because the Accessibility Service cannot reliably read exact platform-reported video duration across all target applications, dwell time is treated as a behavioral proxy. **Swipe Count** is retained only as a supporting interaction metric that helps show whether a long session consists of rapid chaining or slower browsing.

Session Duration is operationalized as continuous active time on a target application, reflecting the prolonged-engagement dimension reported in doomscrolling-related literature. Together, dwell time and session duration describe how long and how continuously the user

remains engaged. They are retained as observable engagement proxies, not as validated biomarkers of doomscrolling.

Negative Sentiment Density (NSD) captures the negative-exposure dimension of doomscrolling through a hybrid operationalization. When usable text exists, analyzable caption and visible-comment units are scored through the text path. When no usable text exists, the item is routed through the VLM fallback and contributes one resolved no-text label to the same session-level exposure estimate. High-OOV text is not rerouted into the VLM; if negativity still cannot be resolved reliably, the session enters the 2-input behavioral fallback and is marked sentiment-unreliable for the main sentiment-based analyses.

NSD is therefore an exposure proxy, not a direct measure of distress. It cannot distinguish crisis-related negativity from ordinary negative commentary or represent the full multi-modal meaning of viewed media.

Session Boundary Definition: A session begins when a target app enters the foreground (detected via `TYPE_WINDOW_STATE_CHANGED`) and ends when: (a) the target app leaves the foreground for longer than a 30-second bridge window, (b) the screen turns off, or (c) the Accessibility Service is terminated. If the user briefly switches away and returns within 30 seconds, the gap is excluded from duration and dwell-time accumulation but the session continues. Recent smartphone-logging studies that reconstruct sessions from app-event streams continue to use 45-second inter-app gaps when direct lock/unlock boundaries are unavailable (Ahmed & Ahmed, 2023; Chen et al., 2023). The present system therefore treats 30 seconds as a stricter study-defined session-continuity heuristic within that recent measurement range so only very brief app switches are merged; longer away intervals are treated as new sessions. The post-deployment sensitivity analysis examines ± 15 -second perturbations alongside the other fixed study priors.

Idle Timeout Heuristic (Active-vs-Idle Disambiguation): To prevent false inflation of dwell time when the device is unattended (e.g., phone left on a looping video), the dwell

accumulator is treated as **active** only when periodic micro-interactions are observed (e.g., TYPE_VIEW_CLICKED, TYPE_TOUCH_INTERACTION_START/END, or equivalent interaction signals exposed by the Accessibility layer). If no micro-interaction occurs within a rolling 45-second window, the ongoing dwell interval is paused and marked as idle exposure until interaction resumes or the next swipe occurs. Recent smartphone-session reconstruction studies continue to treat gaps of about 45 seconds as a plausible continuity delimiter when direct boundary data are unavailable (Ahmed & Ahmed, 2023; Chen et al., 2023). The dwell tracker uses the same order-of-magnitude value as a study-defined inactivity cap to distinguish active viewing from unattended looping: if interaction signals disappear beyond that range, accumulation pauses until interaction resumes or the next swipe occurs.

$$\text{NSD} = \left(\frac{\text{Negative resolvable units}}{\text{Resolvable units in session}} \right) \times 100$$

To formally represent this calculation in mathematical notation:

$$\text{NSD} = \frac{\sum_{i=1}^{N_R} \mathbb{I}(y_i = 1)}{N_R} \times 100$$

Where N_R is the number of session units whose negativity can be resolved either from analyzable text or from the no-text VLM path, and y_i is the binary negative label assigned to resolved unit i . For text-resolved units, negativity follows the VADER criterion $C_i < -0.05$. For no-text VLM-resolved units, negativity is assigned when the returned visual label is SEVERE_NEG or MILD_NEG. The indicator function $\mathbb{I}$ assigns 1 only to resolved units classified as negative and 0 otherwise. This makes the reported NSD a conservative negative-item density rather than a continuous magnitude-weighted affect score.

The -0.05 threshold follows VADER’s standard classification convention, where compound scores below -0.05 are treated as negative sentiment (Hutto & Gilbert, 2014). Caption/comment text units excluded by the language-reliability screen are removed from the text-side evaluation. No-text items routed to the VLM contribute one resolved item-level label each: SEVERE_NEG and MILD_NEG count as negative exposure, while NEUTRAL,

MILD_POS, and SEVERE_POS count as resolved non-negative exposure. High-OOV text is not rerouted into the VLM path. Sessions with very few resolvable units are flagged descriptively during analysis.

Alternative candidates such as swipe-back frequency, time-of-day patterns, and biometric signals were excluded because they either lack enough empirical support in the doomscrolling literature or are technically infeasible within the intended Android monitoring environment. Likewise, although the **internal psychological Antecedents** of the Feedback Loop Model (e.g., anxiety, FoMO) remain conceptually important, they are not measured directly because real-time capture would require intrusive experience sampling or biometric sensors. The system therefore estimates doomscrolling-related risk from one exposure proxy, NSD, together with two behavioral proxies, Video Dwell Time and Session Duration, while swipe or scroll events remain transition signals and supporting logs rather than a separate estimator input.

3.3.8. Analytic Scoring Framework

Week-level DSI is computed on one fixed scoring framework across both study weeks so the analytic scale remains constant. Swipe events serve only as transition markers for dwell-time computation, and prompt-response logs are reported descriptively because they exist only in Week 2. Live prompting may personalize Session Duration and NSD memberships after Week 1, but analytic DSI scoring remains fixed across both weeks.

3.4. Hardware and Software Specifications

Development uses Android Studio (Ladybug 2024.2.1+) with Kotlin 1.9+ on a machine with at least 8 GB RAM. The deployable application targets Android 8.0 (API 26) or higher, while the screenshot-assisted no-text VLM path requires Android 11 (API 30) or higher because it uses `AccessibilityService.takeScreenshot`. Local storage uses

Room (SQLite), sentiment analysis uses a VADER-compatible lexicon-and-rule engine re-implemented in Kotlin for on-device Android deployment, dashboard visualization uses MPAndroidChart, screen-content extraction uses `androidx.accessibilityservice`, and no-text routing uses the Android Accessibility Service together with the on-device Moon-dream 0.5B VLM/VQA module. A physical Android device is used for field testing.

3.5. Method in Developing Software

3.5.1. Methodology Procedure

3.5.1.1. Agile Scrum

Development follows **Agile Scrum** in four iterative one-week sprints covering: (1) requirements finalization and planning, (2) accessibility extraction and local logging, (3) estimator and prompting integration, and (4) testing, refinement, and deployment preparation. Scrum is selected because the Accessibility Service extraction logic requires frequent adaptation as target-app UI structures change, and the fuzzy-logic rule base requires iterative tuning.

3.5.2. Algorithm

The core logic of the system combines **Threshold-Based Rules**, **VADER (Valence Aware Dictionary and sEntiment Reasoner)**, the no-text VLM fallback, and **Fuzzy Logic**. Together, these components govern live prompting, session-level scoring, and week-level DSI computation.

1. Threshold-Based Heuristics: For **live prompting**, the system first applies duration-based gates before deeper inference. Recent JITAI and infinite-scroll studies support burden-aware timing and treat sub-10-minute use as relatively brief while placing sustained scrolling in the 10 to 20+ minute range (Fiedler et al., 2024; Genugten et al., 2025; Hsu et al., 2025; Meinhardt et al., 2025; Rixen et al., 2023; Teepe et al., 2021; Terzimehić &

Aragon-Hahner, 2022; Wang et al., 2023). Accordingly, the system uses 15 minutes as the study-defined live-prompt gate and cooldown interval.

- **Session Duration Gate:** deeper prompting begins after 15 minutes of continuous use.
- **Intervention Cooldown:** after a prompt is shown, the system requires another 15 minutes of qualifying use before a new prompt can be triggered.

Mean Video Dwell Time still contributes to the fuzzy engine once prompt eligibility is met, but it is not implemented as a separate hard pre-inference gate.

Logged session summaries are retained whether or not a live prompt occurs.

2. VADER Sentiment Analysis: VADER is a lexicon- and rule-based sentiment analysis tool designed for social-media text and suited to on-device use (Hutto & Gilbert, 2014).

- **Input:** Extracted caption or comment text (e.g., “This is heartbreaking 🥹”).
- **Process:** Tokenization and lexicon lookup.
- **Output:** A Compound Score ranging from -1.0 (Most Negative) to $+1.0$ (Most Positive). Mathematically, VADER computes this bounded score via the normalization scalar:

$$C = \frac{x}{\sqrt{x^2 + \alpha}}$$

Where x is the sum of the valence scores of recognized words, and α is the standard normalization constant ($\alpha = 15$).

The implementation also retains VADER’s native UTF-8 emoji valence parsing, so emoji-dominant comments (e.g., “🥹🥹” or “💀💀💀”) still contribute affective signal even when lexical Taglish syntax is highly out-of-vocabulary.

Filipino Lexicon Extension: To handle Filipino and Taglish expressions, the system extends VADER with a limited **Minimum Viable Lexicon** (MVL). Recent low-resource sentiment work and lexicon-curation studies support this approach for code-mixed,

resource-constrained settings (Mohammed & Prasad, 2023; Nazir et al., 2026; Tho et al., 2021; Wijayanti & Arisal, 2021).

1. **Corpus Collection:** Before participant deployment, up to 50 high-frequency non-English tokens are compiled from exploratory extraction logs produced during researcher-controlled test sessions or earliest non-study dry runs. The 50-term cap is a study-defined feasibility rule that keeps the lexicon focused on dominant recurrent tokens.
2. **Expert Polarity Rating:** The collected tokens are presented to a panel of three Filipino language teachers and/or linguistics students, who independently rate each token on a sentiment polarity scale from -4 (most negative) to $+4$ (most positive), following VADER’s valence-rating convention.
3. **Inter-Rater Agreement Screening:** A token is retained only if all 3 raters assign values within ± 1 of the mean. Recent sentiment-annotation work shows that agreement can be difficult even with explicit instructions, so the ± 1 screen is used as a conservative agreement rule for retaining only tokens with comparatively stable local judgments before lexicon integration (Krušić, 2024; Äyräväinen et al., 2025).
4. **Integration:** The retained tokens and their mean valence scores are appended to the VADER lexicon dictionary at runtime, enabling the sentiment engine to recognize common Filipino expressions such as “nakakalungkot” (sad), “nakakagalit” (infuriating), or “sana all” (envious aspiration).

This approach keeps the vocabulary focused on recurring tokens while preserving annotation quality and computational efficiency for edge deployment.

3. Taglish Syntax Handling and Fallback Logic: The system uses a lightweight lexicon-based approach to keep inference on-device under mobile resource and privacy constraints. Local and low-resource studies support this choice for Philippine social-media text and

code-mixed sentiment analysis (Co et al., 2022; Cruz & Balahadia, 2022; Mohammed & Prasad, 2023; Nazir et al., 2026; Pacol & Palaoag, 2021).

To address negations such as “hindi maganda” (“not good”), the system implements a localized negation heuristic. A predefined list of common Filipino negation markers (e.g., “hindi”, “di”, “wala”, “wag”) is integrated into VADER’s pre-processing step. If a negation marker is detected within three tokens preceding a recognized sentiment token, the polarity of that token is inverted before the final compound score is computed.

Informal social-media text also contains slang, textual variation, and deliberate misspellings. Unmatched variants remain unmatched and contribute to the OOV calculation rather than being force-mapped into a sentiment label through fuzzy matching.

Dialect-heavy text and purely visual items can still fall outside the reliable text path. The system addresses that case through the no-text VLM fallback:

3.5.3. No-Text Sentiment Fallback via Moondream 0.5B

Short-form video environments can include purely visual items without accompanying text payloads. For that case, the system uses a **Per-Video Routing Heuristic** that calls a VLM fallback built on Moondream 0.5B. Moondream 0.5B is used as a small deployment-fit VLM for on-device edge inference (Lee et al., 2024; Sharshar et al., 2025).

- **Dynamic Routing Logic:** When transitioning to a new item in the feed, if the Accessibility Service successfully extracts usable text payload (caption or visible comments), the item is evaluated via VADER and routed into the text calculation tier. If no usable text is present (e.g., purely visual content without captions), the item enters the Accessibility Service screenshot plus Moondream 0.5B path. High-OOV text is not rerouted into the VLM path; it is handled through the same reliability safeguard that triggers the fallback mode.

- **Data Extraction via AccessibilityService.takeScreenshot:** Visually routed items use an on-demand screenshot after the new item has settled. This is a resource-aware compromise supported by related affective-video and privacy-compliant video analysis work showing that sparse visual evidence can support practical inference under resource constraints (Augusma et al., 2023; Qu et al., 2025; Sharma et al., 2023). To maintain the edge-computing privacy boundary, the captured frame exists only in volatile RAM for VLM inference, ensuring no user screen data is committed to persistent storage.
- **VQA Zero-Shot Computation:** The transient screenshot is processed entirely on-device by the Moondream 0.5B VLM/VQA module. The system uses constrained Visual Question Answering (VQA), and the no-text item is assigned one of five labels: SEVERE_NEG, MILD_NEG, NEUTRAL, MILD_POS, or SEVERE_POS. If the label is SEVERE_NEG or MILD_NEG, the item contributes one negative resolved unit to NSD; otherwise it contributes one resolved non-negative unit.

This route extends NSD coverage to no-text items while remaining within the device’s privacy boundary. It remains an item-level visual classification path rather than a full multimodal interpretation of viewed media (Cortiñas-Lorenzo & Lacey, 2024; Das & Singh, 2023; Johnson et al., 2025; Sharshar et al., 2025; Wei et al., 2021; Zhang & Xu, 2023).

Safety Lock: If the no-text VLM path fails because AccessibilityService.takeScreenshot is unavailable, a screenshot cannot be captured, or VLM inference does not produce a stable item label, the system drops to 2-input fuzzy inference (Video Dwell Time + Session Duration only). Sessions dominated by high-OOV text follow the same fallback path. In either case, unresolved sessions are marked **Sentiment-Unreliable** and excluded from main Doomscroll Severity Index (DSI) statis-

tical evaluation. The degraded rule base and prompt-mapping thresholds are detailed in Appendix B.

Operationally, each viewed item stays on the text path when usable text exists and moves to the VLM fallback only when no usable text exists. If the session-level OOV ratio reaches 50% or more, or if no-text resolution fails, the session is marked **Sentiment-Unreliable**: the sentiment input is dropped, inference proceeds on Video Dwell Time and Session Duration only, Level 3 prompts are disabled, and the session is excluded from main DSI computation. The 50% OOV screen functions as a study-defined reliability rule, and the study reports the number and proportion of sentiment-unreliable sessions to show DSI coverage.

4. Fuzzy Logic Inference System: The system uses fuzzy logic because the target construct is gradual, the rule base must remain inspectable for SME review, and labeled session data are not available for supervised on-device modeling (Pickering et al., 2025; Vashishtha et al., 2023). The fuzzy engine operates at the per-session inference level, while the statistical tests in Chapter 4 are used later for sample-level evaluation. The membership boundaries below define the fixed scoring framework for week-level DSI across both study weeks. In Week 2 live prompting, Session Duration and NSD memberships may be personalized from Week 1 sentiment-reliable sessions, while Video Dwell Time remains fixed (Ikegaya et al., 2025).

- **Fuzzification:** Crisp inputs (Video Dwell Time in s, Negative Sentiment Density in %, Session Duration in min) are converted into fuzzy sets using triangular membership functions. Mathematically, a triangular membership function with bounds (a, b, c) is defined as:

$$\mu_A(x) = \begin{cases} 0 & \text{if } x \leq a \\ \frac{x-a}{b-a} & \text{if } a < x < b \\ \frac{c-x}{c-b} & \text{if } b \leq x < c \\ 0 & \text{if } x \geq c \end{cases}$$

With three input variables and three linguistic states per variable (Low, Medium, High), the rule base contains $3^3 = 27$ rules. The table below presents the membership boundaries used for analytic scoring:

Table 3. Fuzzy Set Membership Boundaries for Analytic Scoring

Variable	Low	Medium	High
Video Dwell Time (s)	0–5	4–20	15–30+
Negative Sentiment Density (%)	0–33	17–83	67–100
Session Duration (min)	0–10	8–20	15–40+

These boundaries combine recent app-use timing evidence with a regular three-part partition for NSD, where the literature does not provide standard per-session negativity thresholds (Azam et al., 2021; Casalino et al., 2022; Cho et al., 2021; Ismail & Al Thani, 2022; Khairuddin et al., 2021; Meinhardt et al., 2025; Muise et al., 2024; Terzimehić & Aragon-Hahner, 2022; Tian et al., 2021). They serve as the pilot scoring structure for this estimator. Session Duration and Negative Sentiment Density act as the main risk axes, while Video Dwell Time mainly adjusts risk by showing how sustained attention is distributed across viewed items.

The table reports support intervals rather than every triangular parameter. In implementation, each Low set is a left-shoulder triangle anchored at zero, each Medium set peaks at the midpoint of its displayed support interval (12 s, 50%, and 15 min), and each High set is a right-shoulder triangle that saturates at the upper cap (30 s, 100%, and 40 min). This setup was chosen because it is the simplest piecewise-linear form that matches the declared Low/Medium/High meanings, keeps adjacent overlap without gaps, prevents overlap between non-adjacent sets, and keeps every parameter easy to check against the table.

The boundaries are interpreted as follows:

- **Video Dwell Time:** Low (0–5 s) represents immediate skip or near-skip behavior, while High (15–30+ s) represents clearly sustained attention to a single item.

- **Negative Sentiment Density:** Low, Medium, and High are defined through an evenly distributed three-term partition over the 0-100% NSD scale because the literature does not provide standard per-session negativity thresholds.
- **Session Duration:** Low (0–10 min) represents brief checking, while High (15–40+ min) represents clearly extended sessions.

The rule base is arranged so that increasing Session Duration should not lower risk when the other inputs are fixed, and increasing Negative Sentiment Density should not lower risk. Increasing Video Dwell Time generally raises risk because it reflects more sustained attention to a single item, but that effect is weaker and more context-dependent than the duration and negativity axes. The resulting rule surface is therefore mostly monotonic rather than strictly monotonic.

These ranges are retained as the fixed analytic scoring framework for week-level DSI across both study weeks and as the Week 1 live default. After the 7-day baseline window, if at least 10 sentiment-reliable sessions are available, the live prompt engine updates Session Duration and NSD memberships from that participant's Week 1 $\frac{Q_{25}}{Q_{50}}$ quantiles; otherwise the default memberships are retained. The ≥ 10 threshold is used to avoid unstable personalization from sparse baseline data. Video Dwell Time remains fixed in both analytic and live modes.

Sensitivity Analysis of Fuzzy Boundaries: Post-deployment, the study conducts a one-at-a-time local sensitivity analysis of the nine boundary parameters by varying each parameter across a pre-specified plausibility range while the others are held constant. For each perturbation, the analysis reports: (1) mean absolute score change, (2) percentage of sessions reclassified (Safe/Warning/Critical), and (3) which parameters produce the largest output instability under the tested range (Dogan, 2021; Shahari & Rasmani, 2024; Shukla & Muhuri, 2025; Vinogradova-Zinkevič, 2023).

- **Rule Evaluation:** The system evaluates a rule base of 27 rules (3 variables $\times$ 3 levels each = 3^3 combinations). **Safe** is assigned to clearly low-risk states and isolated mild signals. **Warning** is assigned to mixed or moderate states without strong multi-axis convergence. **Critical** is assigned when prolonged use and negative exposure converge, with Video Dwell Time acting mainly as an intensifier.

The rule assignments follow these principles:

- every Low/Medium/High combination is assigned exactly once, so the engine has no undefined corner cases
- no rule is **Critical** when Session Duration and Negative Sentiment Density are both Low
- prolonged duration alone does not produce **Critical** when Negative Sentiment Density remains Low
- High NSD is escalated most strongly when at least one other axis is also non-low, especially Session Duration
- High Dwell Time alone is insufficient for **Critical** and mainly intensifies already concerning states
- **Warning** remains the default mixed class whenever the evidence is concerning but not yet strongly convergent

Read in groups, Rules 1, 2, 4, 10, and 19 are assigned **Safe** because at least two axes remain low. Rules 3, 5-8, 11-14, 16, and 20-22 are assigned **Warning** because they show one elevated axis or mixed evidence without strong duration-negativity convergence. Rules 9, 15, 17-18, and 23-27 are assigned **Critical** because prolonged use and negative exposure converge, or because High NSD is reinforced by High Dwell Time even when the total session is still short.

Rule 15 (Medium Dwell, Medium NSD, High Duration) is classified as **Critical** because the session is already prolonged and both other axes are non-low. Rule 23 (High Dwell,

Medium NSD, Medium Duration) is also classified as **Critical** because sustained attention and non-low negative exposure already converge in a session that is no longer brief. Rule 25 (High Dwell, High NSD, Low Duration) reflects the same logic in concentrated form: even a short session can show concentrated negative immersion when viewed items are mostly negative and each item holds attention for a relatively long time.

Representative rules include:

Table 4. Representative Appendix A Rules (6 of 27)

No.	Video Dwell Time	Negative Sentiment Density	Session Duration	Risk Level
1	Low	Low	Low	Safe
14	Medium	Medium	Medium	Warning
27	High	High	High	Critical
26	High	High	Medium	Critical
9	Low	High	High	Critical
19	High	Low	Low	Safe

The complete 27-rule inference base is presented in Appendix A.

Exception Logic: Most rules are arranged so that higher Session Duration or higher Negative Sentiment Density does not reduce risk, but one deliberate exception is retained for interpretive accuracy. Appendix A Rule 9 (Low Dwell, High NSD, High Session Duration) is classified as **Critical** because low dwell should not automatically downgrade a session that is both prolonged and dominated by negative items; in that edge case, rapid chaining across many negative items is treated as more concerning than reassuring. By contrast, Appendix A Rule 19 (High Dwell, Low NSD, Low Session Duration) is not treated as an exception but as a boundary clarification: sustained attention to a single item, without prolonged session time or negative-dominant exposure, is not enough by itself to indicate doomscrolling. This is why the table treats dwell time as a modifier, not the main driver.

- **Defuzzification:** The aggregated fuzzy output is converted back into a crisp “Risk Score” (0–100) using the Center of Gravity (CoG) method. The three output classes occupy equal thirds of the reporting scale, with category cutoffs at 33.33 and 66.67 and output centers at 16.67, 50.00, and 83.33 for Safe, Warning, and

Critical, respectively. Formally, this computation involves rule activation weight (w_i) following the MIN T-norm, and output defuzzification:

$$w_i = \min(\mu_{\text{Dwell}}(x), \mu_{\text{NSD}}(y), \mu_{\text{Duration}}(z))$$

$$\text{Risk Score} = \frac{\sum_{i=1}^n w_i \cdot c_i}{\sum_{i=1}^n w_i}$$

Where w_i is the activation weight of rule i , and c_i is the crisp output center for the mapped risk class. For each fired rule, the smallest membership value determines the rule's activation weight. This prevents one strong input from fully dominating when the other conditions required by that rule remain weak.

Dynamic Weighting vs. Static Percentages: The adopted MIN-based fuzzy inference does not assign flat percentage weights to individual variables. Each fired rule is limited by its weakest required antecedent, so a variable's influence depends on the full combination of inputs rather than on a fixed linear contribution.

Illustrative Worked Example: Given the input tuple (VideoDwellTime = 45 s, NegativeSentimentDensity = 62%, SessionDuration = 35 min), the system computes Video Dwell Time High membership = 1.0, Negative Sentiment Density Medium membership ≈ 0.64 , and Session Duration High membership = 0.80. The dominant firing pattern is Appendix A Rule 24 (High Dwell Time, Medium NSD, High Duration $\rightarrow$ Critical), so the defuzzified Risk Score is 83.33/100 and the resulting intervention level is Level 3 (Critical).

3.6. Method in Evaluation of the Study

The study uses four complementary evaluation methods: the **ISO/IEC 25010 Software Quality Model** for technical quality assessment, TAM for user acceptance evaluation, an SME panel for expert review, and a baseline-intervention comparison of selected logged metrics and self-reported doomscrolling scores. Results are reported separately as **operational feasibility** (successful deployment completion and absence of critical system

failures), **software acceptability** (ISO/IEC 25010, TAM, and SME results), and **estimator plausibility** (baseline convergent association together with sentiment-reliable coverage).

3.6.1. ISO/IEC 25010 Software Quality Evaluation

The study specifically focuses on four key characteristics of the ISO/IEC 25010 model (International Organization for Standardization, 2023) relevant to privacy-sensitive mobile well-being applications:

1. **Functional Suitability:** Assessing whether the heuristic risk-state estimation and intervention features function as intended and cover the specified requirements.
2. **Performance Efficiency:** Evaluating the application's resource usage through both user survey responses and objective profiling metrics such as CPU usage, RAM consumption, and battery drain observed during controlled testing on a small reference-device set, with field-deployment logs used only as descriptive context.
3. **Usability:** Measuring the ease of use, learnability, and user satisfaction with the interface and intervention mechanisms.
4. **Reliability:** Assessing the system's stability and availability during extended periods of monitoring without crashing or being killed by the Android OS.

3.6.2. Technology Acceptance Model (TAM) Evaluation

To assess user acceptance, the study uses the Technology Acceptance Model (TAM). Recent mHealth acceptability reviews continue to treat TAM as a practical user-acceptance framework, with Perceived Usefulness and Perceived Ease of Use retained as core constructs (Adnan et al., 2025). Two TAM constructs are measured:

1. **Perceived Usefulness (PU):** The degree to which users believe the system enhances their ability to manage doomscrolling behavior.

2. **Perceived Ease of Use (PEOU):** The degree to which users believe the system is free of effort to learn and operate.

Each construct is measured using 6 Likert-scale items aligned to these two core TAM constructs and adapted to the doomscrolling risk-estimation context. The TAM instrument is administered in the Phase 5 post-usage survey.

3.6.3. Subject Matter Expert (SME) Evaluation

A panel of **three (3) subject matter experts** evaluates the system's technical quality and risk-estimation logic. The SME panel is composed of:

- One (1) expert in software engineering or mobile application development
- One (1) expert in data science, machine learning, or fuzzy logic systems
- One (1) expert in digital well-being, behavioral psychology, or educational technology

Each SME evaluates the system using a structured rubric covering: (1) the technical soundness of the hybrid risk-estimation framework, including VADER integration and fuzzy inference behavior, (2) the plausibility of the selected input ranges for Video Dwell Time, Negative Sentiment Density, and Session Duration, (3) the internal coherence of the 27-rule base, including whether nearby rules change in a sensible direction across duration and negative exposure and whether exception rules such as rapid negative-content chaining are reasonable, (4) the quality of the system architecture and privacy implementation, (5) the appropriateness of the intervention design, and (6) overall software quality using the same ISO/IEC 25010 characteristics evaluated by end users. SME responses are collected using a 5-point Likert scale and analyzed descriptively (mean, standard deviation) alongside qualitative feedback. This SME review serves as narrative expert appraisal of the rule logic and as an expert plausibility check, not as formal empirical calibration or a content-validity index.

3.6.4. Testing Method

The study uses both **black box** and **white box** testing. Black-box testing checks whether visible functions behave as expected under realistic use, while white-box testing checks whether the internal scoring, routing, and data-flow logic behave correctly at code and module level.

3.6.4.1. Black Box Testing Procedure

The black-box procedure evaluates the application through expected inputs, user actions, and observed outputs without relying on internal code inspection. The full application is deployed on a small set of physical devices representing the screen sizes and Android versions available to the researchers to verify that usage patterns (e.g., prolonged session with high negative content) produce the expected risk classification and intervention, including correct routing of no-text cases into the VLM fallback and correct degradation to the 2-input fallback when sentiment remains unresolved. The same procedure checks user-visible flows such as onboarding, permission handling, dashboard display, logging continuity, and prompt presentation.

3.6.4.2. White Box Testing Procedure

The white-box procedure evaluates the internal logic and data flow of the system. The Kotlin-implemented VADER-compatible scoring engine and fuzzy-logic rules are tested in isolation using JUnit to verify that known inputs produce expected outputs. The Accessibility Service-to-database pipeline is also tested to confirm correct data flow and that interventions trigger from logged state changes. This same layer verifies text-path routing, no-text VLM routing, and expected fallback to the 2-input safety mode when sentiment remains unresolved.

3.7. Data Collection Scheduling

Data collection follows five sequential phases. Phase 1 covers recruitment, screening, consent, and group assignment. Phase 2 covers baseline profile completion and application setup. Phase 3 covers the one-week baseline logging period with prompts disabled. Phase 4 covers the one-week deployment period in which only the intervention group receives adaptive prompts while the control group continues logging-only. Phase 5 covers the post-usage survey and final retrieval of aggregate local logs.

3.8. Data Gathering

Data gathering is conducted through five phases:

Phase 1: Recruitment, Screening, and Consent Potential participants are recruited through non-probability purposive-convenience sampling using accessible channels such as academic networks, peer referrals, and online groups. Interested individuals undergo screening to confirm eligibility based on age, active use of at least one target short-form video platform, and study-device compatibility, including whether the phone can install the study build, maintain the required permissions, and keep the monitoring service active under ordinary use. Eligible participants are then provided with an informed consent form explaining the study's purpose, procedures, data to be collected, risks, safeguards, and right to withdraw. After consent and eligibility confirmation, participants are assigned to either the intervention group or the logging-only control group using the concealed permuted-block allocation list described earlier, with an initial 1:1 allocation target of 25 per group. Both groups receive the same baseline setup and Week 1 logging procedure; the distinction becomes operational only in Week 2 when prompts are enabled for the intervention group.

Phase 2: Baseline Profile Form and Application Setup Prior to deployment, each participant completes a baseline profile form capturing demographic information, estimated short-form video use, and general platform-use habits. The participant is then assisted in installing the application and enabling the necessary Android permissions. Each participant

is assigned a unique study code used to link logs and survey forms; the name-to-code list is retained separately by the researchers. This form provides contextual information but is not treated as the primary source of behavioral measurement.

Week 1 provides the fixed-prior baseline DSI and self-report anchor. At the end of the 7-day baseline, the live prompt engine may lock participant-specific Session Duration and NSD memberships from Week 1 sentiment-reliable sessions while leaving Video Dwell Time fixed. This follows recent JITAI evidence favoring personalized intervention criteria over uniform ones for live prompting (Ikegaya et al., 2025). If reliable baseline coverage is insufficient, the default priors are retained.

Phase 3: Week 1 Baseline Logging Phase During Week 1, the app’s built-in logger records quantitative usage metrics locally on the participant’s device while adaptive intervention prompts remain inactive. This baseline period establishes the participant’s observed usage pattern under the study conditions and provides the baseline DSI and self-report anchor used for convergent association testing.

The Week 1 logged data include:

- **Usage Metrics:** Session duration, video dwell time, and supporting interaction counts such as swipe totals.
- **Sentiment-Related Indicators:** Session-level sentiment scores and Negative Sentiment Density estimates.
- **Reliability Events:** Application crashes, service interruptions, and related system-stability notes.
- **Performance Indicators:** CPU, RAM, and battery observations gathered through limited in-app logging and controlled profiling on a small reference-device set, with heterogeneous participant-device traces used descriptively.

At the conclusion of Week 1, participants complete the short-form 4-item **Doomscrolling Scale** (Satici et al., 2023; Sharma et al., 2022), using Items 1, 2, 10, and 12 of the original

instrument and a 7-day recall instruction stem (e.g., “In the past week...”; see Appendix D), to provide the baseline self-report anchor for convergent association. The original instrument was validated as a general self-report scale, so changes between administrations are interpreted as short-horizon self-report differences.

Phase 4: Week 2 Deployment Phase During Week 2, **intervention-group** participants receive adaptive digital mindfulness prompts through a live prompt engine that retains the fixed rule base and fixed Video Dwell Time memberships but, when available, replaces the global Session Duration and NSD memberships with participant-specific Week 1 quantile-derived bounds. **Control-group** participants continue logging-only operation with prompts disabled via a configuration flag. Both groups continue to have usage and sentiment-related metrics logged, while prompt-response data for the intervention group are recorded descriptively.

For behavioral comparison, raw elapsed session duration and raw elapsed video dwell time are treated as the primary time-based metrics. **Prompt-excluded active-use metrics** are also retained: when a prompt appears, the active content-view interval is closed, and time spent inside the system-generated intervention is excluded until the participant returns to the target platform. These prompt-excluded metrics are reported as supplementary traces of prompt-interrupted use. Observed Week 2 differences are interpreted as short-term behavioral differences under the study conditions.

Potential confounds include time-of-day and day-of-week variation, external events during the two-week window, and device-specific differences in processing speed or background-process management. The randomized two-group design helps balance these factors across study arms, while the within-subject Week 1 to Week 2 comparison in the intervention group accounts for stable individual differences.

At the conclusion of Week 2, participants complete the same short-form 4-item **Doomscrolling Scale** under the same 7-day recall instruction stem (Appendix D) to provide the post-intervention self-report anchor for paired comparison.

Phase 5: Post-Usage Survey and Final Data Retrieval To align data structures for baseline convergent association and descriptive week-level comparison, logged session summaries over each 7-day period are aggregated into a composite **Doomscroll Severity Index** per user per week. This analytic DSI uses the same fixed-prior scoring framework across both study weeks. Unlike the live Week 2 prompt engine, the DSI engine is not personalized after Week 1, so week-level summaries remain on one fixed scale:

$$DSI_w = \left(\frac{1}{m_w} \right) \sum_{s=1}^{m_w} RiskScore_s$$

where m_w is the number of week- w sessions classified as sentiment-reliable and therefore doomscrolling-evaluable. Sessions flagged as **Sentiment-Unreliable** are excluded from DSI computation. DSI is a mean of retained session scores rather than a cumulative sum, so these exclusions affect coverage more than scale. Week-specific DSI values are retained for descriptive comparison, convergent testing is anchored to the baseline week only, and DSI coverage is reported by week and study arm in terms of retained participants, retained sessions, and excluded sentiment-unreliable sessions.

Interpretation of Classification Metrics: Traditional supervised classification metrics (e.g., precision, recall, F1, AUROC) are not estimated because session-level ground-truth labeling would require interrupting active scrolling. The study therefore uses convergent association as the main plausibility check and interprets the system as a heuristic behavioral risk estimator.

After the deployment period, participants complete a structured post-usage survey instrument with three components: (a) the ISO/IEC 25010 software quality evaluation, (b) the TAM assessment, and (c) open-ended feedback items. At the same stage, aggregate local

logs are exported using the participant study code only and transferred to a password-protected research folder accessible only to the researchers.

For the ISO/IEC 25010 component, the survey uses a 5-point Likert Scale to evaluate the four target quality characteristics: Functional Suitability, Performance Efficiency, Usability, and Reliability. For the Usability sub-characteristic, all 10 items of the System Usability Scale (SUS) are adopted verbatim, and the standard SUS scoring formula is applied to yield a composite usability score on a 0–100 scale. The SUS is a standardized usability measurement tool whose validity for digital health applications has been supported by a meta-analysis of 114 apps yielding a benchmark mean of 68.05 (SD 14.05) (Hyzy et al., 2022). For the remaining three sub-characteristics (Functional Suitability, Performance Efficiency, and Reliability), the study uses researcher-developed items derived from the ISO/IEC 25010 sub-characteristic definitions (International Organization for Standardization, 2023) and refined through pre-deployment expert review.

For the TAM component, 12 Likert-scale items (6 for Perceived Usefulness, 6 for Perceived Ease of Use) aligned to the study’s contextualized TAM constructs are included to assess user acceptance of the system.

3.8.1. Instrument Review and Reliability Safeguards

Before deployment, the combined survey instrument undergoes structured expert content review to check item relevance, clarity, redundancy, and alignment with the study constructs. Reviewer comments are consolidated, and unclear or overlapping items are revised before administration.

After data collection, internal consistency for each multi-item sub-scale is estimated through Cronbach’s alpha and interpreted as an index of item coherence, with appropriate caution given the pilot sample size and the mix of standardized and adapted items.

Open-ended feedback items are used as descriptive support for the survey findings. Two researchers independently review the responses, reconcile coding differences through discussion, and report recurring issue categories.

The alignment between the study's research variables, data sources, and instruments is summarized in the following table.

Table 5. Variable–Instrument Alignment Matrix

Research Variable	Data Source	Instrument / Tool	Phase
Baseline usage profile	Self-report / Logs	Baseline profile form + Week 1 logs	Phase 2 & 3
Session Duration	Automated system logs	Built-in application logger	Phase 3 & 4
Video Dwell Time (s)	Automated system logs	Built-in application logger	Phase 3 & 4
Negative Sentiment Density (%)	Automated system logs	VADER text-path + VLM no-text path + Built-in logger	Phase 3 & 4
Prompt-response profile (descriptive)	Automated system logs	Built-in application logger	Phase 4
Pre/Post Doomscrolling Scale Score	Self-report	Short-form 4-item Doomscrolling Scale (Appendix D)	End of Phase 3 & 4
Baseline convergent association	Week 1 DSI + Week 1 Doomscrolling Scale	Bivariate Correlation (Pearson's r / Spearman's ρ)	End of Phase 3
Functional Suitability	Likert-scale survey	Researcher-developed ISO/IEC 25010 items + expert review	Phase 5
Performance Efficiency (subjective)	Likert-scale survey	Researcher-developed ISO/IEC 25010 items + expert review	Phase 5
Performance Efficiency support (objective profiling)	Profiling records / logs	Controlled profiling + field-deployment metrics	Phase 3 & 4

Table 5. Variable–Instrument Alignment Matrix (continued)

Research Variable	Data Source	Instrument / Tool	Phase
Usability	Likert-scale survey	System Usability Scale (SUS), benchmarked in Hyzy et al. (2022)	Phase 5
Reliability	Likert-scale survey	Researcher-developed ISO/IEC 25010 items + expert review	Phase 5
Perceived Usefulness (TAM)	Likert-scale survey	Contextualized TAM PU items	Phase 5
Perceived Ease of Use (TAM)	Likert-scale survey	Contextualized TAM PEOU items	Phase 5
Expert quality assessment	SME evaluation	Structured rubric (ISO 25010 + risk-estimation logic)	SME

3.9. Respondents of the Study and Sampling Technique

The respondents of this study are Filipino Android users aged 18 years and above who actively use at least one of the target short-form video platforms. The findings apply to this target population and to the non-probability nature of the sample. The sampling technique employed is **purposive-convenience sampling** because participants must meet specific technical and behavioral eligibility criteria and be reachable through channels accessible to the researchers.

The study reports how many participants used TikTok, Facebook Reels, and Instagram Reels during the field evaluation, including platform mix by study arm. Per-platform extraction success rates and analyzable-session counts are also reported, and cross-platform interpretation is limited to platforms that yielded stable extraction throughout the deployment window.

Inclusion Criteria:

- Must be aged 18 years or older to provide independent informed consent.
- Must own an Android smartphone (Version 8.0 or higher) that passes study-compatibility screening for installation, permissions, and background-service stability.
- Must be an active user of at least one target short-form video platform (TikTok, Facebook Reels, or Instagram Reels).
- Must consent to the installation of the application and the collection of usage data.

Sample Size Justification: A minimum of $N = 50$ participants is targeted, split as evenly as possible into an intervention group and a control group, with an initial allocation target of $n = 25$ per group. This target supports a balanced pilot deployment and post-usage evaluation within an undergraduate thesis scope; it is not intended as a fully powered efficacy sample. Analyses that exclude sentiment-unreliable sessions or participants with insufficient baseline data report their effective sample sizes separately. With a minimum sample of about

50 participants, between-group estimates are interpreted with emphasis on effect sizes and 95% confidence intervals alongside p-values.

Selection Bias Acknowledgment: Purposive-convenience sampling may overrepresent users who are already receptive to self-regulation tools. This self-selection tendency can affect perceived acceptance and compliance estimates.

3.10. Statistical Treatment

Research Question 1 is addressed descriptively through the design-and-development sections earlier in this chapter and is not treated as an inferential hypothesis test. The evaluation data are analyzed using the following statistical methods, each selected for a specific data source and objective.

1. Descriptive Summaries (Survey Data — ISO 25010 & TAM): The Likert-scale responses from the ISO/IEC 25010 software quality evaluation survey and the TAM instrument (Perceived Usefulness and Perceived Ease of Use subscales) are summarized using **mean scores** and standard deviation.

Mean scores are described on a 5-point Likert scale, where 3.00 represents the neutral midpoint. For evaluative decisions, the study uses pre-set favorable thresholds of ≥ 3.50 for ISO/IEC 25010 and TAM constructs and ≥ 4.00 for SME ratings. These are study-defined favorable benchmarks rather than universal cutoffs.

2. Objective Acceptance Thresholds for Evaluation: To reduce post-hoc rationalization of evaluation results, the following pre-set interpretation thresholds are established as decision aids for Chapter 4 rather than as stand-alone proof of software success:

- **System Usability Scale (SUS):** A composite score of ≥ 70 is treated as the study's usability target, set slightly above the 68.05 digital-health benchmark mean reported by Hyzy et al. (2022).

- **ISO/IEC 25010 Subscales:** A mean score of ≥ 3.50 out of 5 is treated as the study-defined favorable-evaluation target for Functional Suitability, Performance Efficiency, and Reliability.
- **TAM Constructs:** A mean score of ≥ 3.50 is treated as the study-defined favorable-evaluation target for Perceived Usefulness and Perceived Ease of Use.
- **SME Evaluation:** A mean rating of ≥ 4.00 out of 5 is treated as the study-defined favorable-evaluation target for expert review.

Objective CPU, RAM, and battery observations are summarized descriptively across the deployment and controlled profiling sessions. Given device differences in hardware and manufacturer background-process behavior, these profiling results are interpreted as supporting evidence for Performance Efficiency instead of as a uniform decision criterion. Results below these thresholds are reported as unmet software-acceptability targets in Chapter 4. Operational feasibility and estimator plausibility are reported separately. Given the pilot-sized sample and the possibility of analyzable-sample reduction from sentiment-unreliable sessions, inferential results are interpreted primarily through effect sizes, confidence intervals, and coverage rather than pass-fail significance alone.

3. Cronbach's Alpha (Instrument Reliability): To assess internal consistency, **Cronbach's alpha** (α) is computed for each sub-scale (Functional Suitability, Performance Efficiency, Usability, Reliability, PU, PEOU). A pre-set coefficient threshold of $\alpha \geq 0.70$ is adopted as a conventional descriptive benchmark for this study. Given the pilot sample and the mix of standardized and adapted items, alpha is treated as a reliability check rather than as a stand-alone validation claim.

4. Primary Behavioral Analysis: The main quantitative comparison for Research Question 2 uses an **Independent Samples t-test on change scores** (Week 2 minus Week 1) to compare observed short-term behavioral differences between the intervention group and the control group. Using change scores rather than raw Week 2 values adjusts for chance

baseline imbalance and individual differences. As a secondary within-subject analysis, a **Paired Samples t-test** compares Week 1 and Week 2 within the intervention group. Given the pilot sample, these tests are treated as exploratory analyses that estimate short-term observed differences under the study conditions rather than precise intervention-effect magnitudes. Separate primary tests are conducted for:

- Mean daily **raw elapsed** session duration (minutes)
- Mean daily **raw elapsed** video dwell time (s)
- Mean daily Negative Sentiment Density (%)
- Pre/Post Doomscrolling Scale Score (Self-Reported)

For the Negative Sentiment Density outcome, mean daily NSD is computed only from sentiment-reliable sessions. Days or participant-weeks with no sentiment-reliable sessions contribute missing data for that specific outcome and are excluded only from that specific NSD comparison, with the analyzed n reported.

As a supplementary transparency check, the same tests are also conducted on **prompt-excluded active-use metrics** for session duration and video dwell time. The difference between raw elapsed and prompt-excluded results is reported to show the measurement-rule effect.

To control family-wise Type I error across the four change-score comparisons, p-values are adjusted using the **Holm-Bonferroni sequential correction** at an initial family-wise $\alpha = 0.05$. For each comparison, **Cohen's d** and its 95% confidence interval are reported alongside the Holm-adjusted p-value to support interpretation of effect size, not statistical significance alone.

Handling of Assumption Violations: For the primary independent-samples t-test, normality of change scores within each group is checked with Shapiro-Wilk and equality of variances with Levene's test. If normality is severely violated, the **Mann-Whitney U test** is

used as fallback. If variances are unequal but normality holds, Welch's t-test is applied. For the secondary paired t-test, normality of within-subject differences is checked with Shapiro-Wilk; if violated, the **Wilcoxon Signed-Rank Test** is used. If analyzable sample sizes become too small for stable inferential interpretation after outcome-specific exclusions, the affected comparison is reported descriptively with its effective n and directional pattern rather than being overstated as confirmatory evidence.

5. Descriptive Analysis of Prompt-Response Logs: Week 2 prompt-response data (e.g., continue, dismiss, take-break) are summarized using frequency counts and percentages. These logs are treated as descriptive evidence of intervention interaction only and are not analyzed as a baseline-comparable outcome variable.

6. Baseline Convergent Association Analysis: Baseline convergent association is assessed on **Week 1 data only**. First, a component-wise correlation table reports the bivariate association between each individual input variable (mean daily session duration, mean daily video dwell time, mean daily NSD) and the baseline Doomscrolling Scale score. Second, the bivariate correlation between the full composite DSI (computed with fixed priors) and the Doomscrolling Scale is computed. The pattern of component-wise versus composite correlations is discussed descriptively to assess whether the multi-variable composite associates more strongly with self-reported doomscrolling than session duration alone. Only participants with at least three sentiment-reliable sessions in Week 1 are included; the effective sample size ($N_{\text{effective}}$) is reported. This coverage rule follows recent digital-measure practice for forming week-level aggregates from repeated observations, with three observations commonly used as a pragmatic weekly minimum while stronger reliability claims require more data (Buekers et al., 2025; Meyer et al., 2022; Ratitch et al., 2023; Yao et al., 2021). Here, the ≥ 3 rule functions as a pragmatic floor for the baseline aggregate. **Pearson's r** is used when linearity and normality assumptions hold; otherwise **Spearman's**

ρ . Results are interpreted as a baseline plausibility check rather than as diagnostic validation.

7. Interpretation: For these exploratory pilot analyses, the nominal level of significance is set at $\alpha = 0.05$. Statistically significant differences between Week 1 and Week 2 are interpreted as observed short-term behavioral differences under the study conditions. The two-group randomized design supports stronger inference than a one-group pre/post comparison, but findings are not interpreted as long-term intervention effects or calibrated validation. Prompt-excluded active-use metrics are interpreted separately because they incorporate the prompt-display rule.

8. Descriptive Coding of Open-Ended Feedback (Phase 5): Open-ended responses from the post-usage survey are analyzed through a light inductive coding procedure. Two researchers read the responses in full, group them into recurring issue categories, and reconcile coding differences through discussion before the categories are summarized. The goal is to identify usability issues, perceived intervention value, and implementation constraints not captured by the Likert-scale measures. This step supports interpretation of the survey results and surfaces recurring concerns for Chapter 4 discussion.

9. Missing Data and Exclusion Criteria: For the convergent correlation analysis, participants need at least three sentiment-reliable sessions in Week 1 ($m_1 \geq 3$) for Week 1 DSI to enter the correlation dataset. This rule includes the edge case $m_1 = 0$, where DSI is undefined by construction. The effective sample size ($N_{\text{effective}}$) after this exclusion is reported explicitly. For the NSD outcome, participants require at least one sentiment-reliable day in each compared week to contribute a valid week-level comparison; otherwise that specific outcome is treated as missing for that participant. For paired pre/post analyses, complete paired observations are required per outcome variable, and the analyzed n for each test is reported.

3.11. Ethical Considerations

This study involves the monitoring of human participants' social media usage behavior, and as such, strict ethical protocols are observed to protect respondent welfare, privacy, and autonomy.

Institutional Review: Prior to data collection, the research protocol, instruments, and informed consent forms are submitted for review and approval by the university's Research Ethics Committee (or Institutional Review Board). No data collection activities commence until formal ethical clearance is obtained.

Informed Consent: All participants are provided with a detailed informed consent form explaining the study's purpose, procedures, data collected, risks, and benefits. Participation is entirely voluntary. The study version requires explicit user activation of the Android Accessibility Service through the device's system settings, providing the study's operational consent layer for text, interaction access, and the service-declared screenshot capability used for no-text visual processing.

Data Protection and Privacy: Consistent with the Privacy by Design principle described in the System Architecture section, the study version processes scraped textual content and temporary no-text screen frames in RAM only, with no persistent storage of raw text data or captured frames. Incidental third-party text, such as usernames or interface labels, is processed briefly and filtered where identifiable so non-participant data is not retained. Only aggregated metrics (e.g., sentiment scores, session duration) are retained. Logs and survey instruments are linked only through a participant study code, while the name-to-code key is stored separately by the researchers. Aggregate exported logs are transferred to a password-protected research folder accessible only to the researchers. The system does not transmit any data to external servers during routine operation.

Right to Withdraw: Participants may withdraw from the study at any time without penalty.

Upon withdrawal, all locally stored usage data is deleted from the participant's device. If withdrawal occurs before final dataset consolidation, any exported participant record linked to that study code is also deleted from the research dataset.

Minimization of Harm: The study does not involve clinical diagnosis or treatment of mental health conditions. The mindfulness interventions (breathing exercises, break reminders) are designed as general wellness prompts, not therapeutic interventions.

The methodological and technical limitations arising from the research design and system architecture of this study are documented in Chapter 1 under **Limitations**.

Chapter 4

PRESENTATION, ANALYSIS AND INTERPRETATION OF DATA

This chapter presents the results, analysis, and interpretation of the simulated REDU export located in `simulated-export/redu-p01-p50-20260523-224114`. The available export contains app-generated session logs, daily summaries, prompt events, reliability events, risk-personalization rows, pre/post Doomscrolling Scale responses, ISO/IEC 25010 survey responses, SUS responses, TAM responses, and coded user open-ended feedback. The SME rating and SME open-ended feedback files are present but blank, so expert-evaluation results are treated as pending rather than estimated or invented. The analysis dataset contains 50 study-coded participants, 2,051 retained sessions, 1,227 daily-summary rows, 1,152 prompt-event rows, 8,549 reliability-event rows, 50 risk-personalization rows, 100 Doomscrolling Scale rows, 50 ISO/IEC 25010 survey rows, 50 SUS response rows, 50 TAM response rows, and 50 coded user feedback rows. All statistics in this chapter were computed from the CSV files in the simulated export.

4.1. Presentation of Results

The results are organized according to the five research questions. Research Questions 1 to 4 are supported by the available logs and user-response files. Research Question 5 remains pending because the SME rating and comment files contain no completed expert entries.

4.1.1. Research Question 1: Privacy-Preserving Mobile Architecture and Estimation Framework

The first research question asks what privacy-preserving mobile architecture and estimation framework can support doomscrolling-related risk estimation on the target short-form video platforms using behavioral indicators and sentiment-related indicators when reliably resolvable, with 2-input behavioral fallback for sentiment-unreliable sessions.

Table 6. System Architecture and Estimation Framework Result Matrix

System Component	Intended Function	Implementation Evidence	Result Status
Data Access Layer	Collect target-app sessions, platform labels, timing, dwell estimates, swipe counts, and reliability events through the Android monitoring pipeline.	The export contains 2,051 sessions across TikTok, Facebook Reels, and Instagram Reels, with corresponding foreground, background, and finalization events.	Met in simulated export
Business Logic Layer	Route analyzable sessions through the sentiment-reliable path, compute NSD when available, and assign risk score and risk level.	The export contains risk scores for all sessions and NSD values for 1,706 sentiment-reliable sessions.	Met for log scoring
2-Input Safety Fallback	Use behavioral inputs when sentiment is unresolved or unreliable and flag affected sessions for exclusion from NSD/DSI analyses.	There were 345 sentiment-unreliable sessions, all paired with HIGH_OOV reliability events and blank NSD values.	Met as fallback flagging
Presentation Layer	Deliver prompts to intervention participants according to risk state while leaving the control group logging-only.	The prompt log contains 576 shown prompts, all from intervention participants; control participants have no prompt events.	Met in exported logs
Privacy Safeguards	Retain aggregate metrics and study codes without persistent raw text or screenshots in the export.	The exported files contain study codes, numeric metrics, platform labels, prompt actions, and reliability codes only.	Met in export contents

The architecture evidence shows that the simulated export followed the planned privacy-preserving logging boundary. The logs contain derived metrics and reliability codes, but

no raw captions, comments, screenshots, names, or free-text user content. The main implementation constraint visible in the export is sentiment coverage: 16.8% of sessions were marked sentiment-unreliable and therefore could not contribute NSD values.

Table 7. Black Box and White Box Testing Summary

Testing Area	Expected Behavior	Result	Remarks
Onboarding and Permissions	The app installs, opens, and guides the participant through required permissions.	Not in export	No onboarding test-result file is included.
Accessibility Extraction	The service detects target app sessions, content transitions, and app-state changes.	Pass in logs	2,051 sessions and matching target foreground/background events were exported.
Text Sentiment Path	The sentiment path scores analyzable content and contributes to NSD.	Pass in logs	1,706 sessions include reliable NSD values.
No-Text / Unreliable Fallback	Unresolved or high-OOV sessions are routed away from main sentiment-based analysis.	Pass in logs	345 sessions were flagged sentiment-unreliable with blank NSD values.
Fuzzy Inference	Each retained session receives a RiskScore and Safe, Warning, or Critical level.	Pass in logs	All 2,051 sessions include risk scores and risk-level labels.
Prompt Triggering	The appropriate intervention prompt appears when risk thresholds are reached.	Pass in logs	576 prompt displays were recorded for intervention participants only.
Local Logging	Aggregate metrics are stored under study codes without raw text or captured frames.	Pass in export	CSV contents are aggregate and study-code based.

The testing summary is based on exported behavior rather than a separate test-report form. It supports the presence of the logging, scoring, fallback, and prompt-event pipeline in the simulated dataset, but it does not replace a formal device-level black-box and white-box test report.

4.1.2. Evaluation Method, Respondents, and Instrument

Table 8. Evaluation Method, Respondents, and Instrument Status

Evaluation Component	Planned Source	Data Available in Export	Status
Behavioral log evaluation	Week 1 and Week 2 app-generated usage records	2,051 sessions, 1,227 daily summaries, 8,549 reliability events, and 1,152 prompt events.	Complete for log analysis
User evaluation	Post-usage ISO/IEC 25010, SUS, and TAM survey instrument	50 ISO/IEC 25010 rows, 50 SUS rows, 50 TAM rows, and 50 coded user feedback rows.	Complete for user evaluation
Subject matter expert evaluation	Expert evaluation form covering technical design, privacy safeguards, heuristic logic, and intervention structure	SME files are present, but rating and comment cells are blank.	Pending
Open-ended feedback	Post-usage comments and expert comments	50 coded user feedback rows; SME feedback blank.	Complete for users; SME pending

The export supports behavioral log analysis, self-report change analysis, user survey scoring, and user feedback coding. Expert scoring and expert qualitative interpretation remain pending.

4.1.3. Likert Scale and Scoring Method

Table 9. Scoring Method and Interpretation Rules

Instrument / Measure	Scoring Basis	Interpretation Range	Status
ISO/IEC 25010 user survey	Mean of 5-point Likert items by quality characteristic.	Favorable target: mean ≥ 3.50 .	Applied
Technology Acceptance Model survey	Mean of 5-point Likert items for PU and PEOU.	Favorable target: mean ≥ 3.50 .	Applied
System Usability Scale	Standard SUS scoring to a 0-100 composite.	Favorable target: score ≥ 70 .	Applied
Subject matter expert evaluation	Mean of 5-point expert-rating items by evaluation area.	Favorable target: mean ≥ 4.00 .	Pending

The interpretation thresholds are retained for consistency with Chapter 3. User thresholds are applied in the following sections, while SME thresholds remain pending until expert ratings are completed.

4.1.4. Respondent Profile and Platform Exposure

Table 10. Participant Profile Summary

Profile Variable	Intervention Group	Control Group	Total	Remarks
Enrolled participants	25	25	50	Balanced by study code
Completed Week 1	25	25	50	At least one Week 1 session
Completed Week 2	25	25	50	At least one Week 2 session
Included in primary behavioral log analysis	25	25	50	Complete week-level log coverage
Included in baseline DSI coverage analysis	25	25	50	All had ≥ 3 reliable Week 1 sessions

All 50 simulated participants were represented in both study weeks. No participant was excluded from the available log-based analyses because of missing week-level session data.

Table 11. Target Platform Use and Extraction Coverage

Platform	Participants Using Platform	Retained Sessions	Extraction Success Rate	Interpretation Notes
TikTok	50 / 100.0%	809	82.6%	Largest session share; 668 reliable sessions.
Facebook Reels	50 / 100.0%	473	86.3%	Fewest sessions but highest reliable-session rate.
Instagram Reels	50 / 100.0%	769	81.9%	Broad coverage with slightly lower reliability.
Multiple-platform use	50 / 100.0%	2,051	83.2%	All participants had sessions on at least two target platforms.

Platform coverage was broad in the simulated export. All target platforms were represented for all participants, and the overall sentiment-reliable session rate was 83.2%. Facebook Reels had the highest extraction success rate, while Instagram Reels had the lowest, although all three platform rates remained above 80%.

4.1.5. Research Question 2: Week 1 to Week 2 Behavioral and Self-Report Changes

The second research question asks what short-term Week 1-to-Week 2 changes are observed in selected logged usage metrics and self-reported doomscrolling scores between the intervention group and the logging-only control group, and within the intervention group across the same period. The simulated export supports both the logged usage metrics and the pre/post Doomscrolling Scale scores.

Table 12. Primary Between-Group Change-Score Comparison

Outcome Variable	Analyzed n	Intervention Change Mean (SD)	Control Change Mean (SD)	Test Used	Effect Size (95% CI)	Holm- Adjusted p
Mean daily raw elapsed session duration (minutes)	I=25; C=25	0.59 (1.59)	-0.38 (2.04)	Welch t	d=0.53 (CI -0.03 to 1.10)	0.200
Mean daily raw elapsed video dwell time (seconds)	I=25; C=25	-0.10 (2.79)	0.34 (2.39)	Welch t	d=-0.17 (CI -0.73 to 0.39)	1.000
Mean daily Negative Sentiment Density (%)	I=25; C=25	-0.70 (5.09)	-1.07 (6.44)	Welch t	d=0.06 (CI -0.49 to 0.62)	1.000
Pre/Post Doomscrolling Scale score	I=25; C=25	-0.46 (0.42)	-0.05 (0.42)	Welch t	d=-0.97 (CI -1.56 to -0.38)	0.005

The log-derived change-score results did not show a statistically reliable between-group difference after Holm adjustment. The strongest observed log pattern was for raw elapsed session duration: the intervention group increased by 0.59 minutes on average, while the control group decreased by 0.38 minutes, but the confidence interval for the effect size crossed zero. In contrast, the self-reported Doomscrolling Scale showed a larger favorable change in the intervention group than in the control group, with an adjusted p-value of 0.005 and a medium-to-large effect-size estimate. This pattern suggests that users reported less doomscrolling after the intervention period even though the short-term behavioral logs did not show a corresponding reliable reduction.

Table 13. Secondary Within-Intervention Pre/Post Comparison

Outcome Variable	Analyzed n	Week 1 Mean (SD)	Week 2 Mean (SD)	Test Used	p-value / Effect Size
Mean daily raw elapsed session duration (minutes)	25	16.74 (11.39)	17.33 (11.43)	Paired t	p=0.074; dz=0.37
Mean daily raw elapsed video dwell time (seconds)	25	13.61 (7.51)	13.50 (7.37)	Paired t	p=0.854; dz=-0.04
Mean daily Negative Sentiment Density (%)	25	42.15 (23.85)	41.45 (24.25)	Paired t	p=0.501; dz=-0.14
Doomscrolling Scale score	25	3.63 (1.82)	3.17 (1.82)	Paired t	p less than 0.001; dz=-1.08

The within-intervention comparison similarly showed no statistically reliable Week 1-to-Week 2 change in the available log metrics. The intervention group's session duration increased slightly, dwell time was essentially unchanged, and NSD decreased slightly. The Doomscrolling Scale, however, decreased from 3.63 to 3.17 within the intervention group. Because the control group is needed to contextualize ordinary week-to-week fluctuation, the within-intervention result should be interpreted together with the primary between-group change-score comparison.

Table 14. Supplementary Prompt-Excluded Active-Use Metrics

Metric	Analyzed n	Raw Elapsed Result	Prompt-Excluded Result	Measurement Interpretation
Session duration	I=25; C=25	Change: I=0.59, C=-0.38 min; d=0.53	Change: I=0.57, C=-0.38 min; d=0.52	Excluding prompt time changed the intervention estimate only slightly.
Video dwell time	N/A	Raw dwell change: I=-0.10, C=0.34 s	Not exported	The export does not contain a separate prompt-excluded dwell column.

Prompt-excluded duration was close to raw elapsed duration, suggesting that prompt-display time did not materially change the simulated session-duration interpretation. Dwell-time adjustment cannot be evaluated from this export because prompt-excluded dwell time was not exported.

4.1.6. Prompt-Response and Reliability Log Summary

Table 15. Prompt-Response Summary for Intervention Participants

Prompt / Response Indicator	Frequency	Percent	Interpretation
Prompts shown	576	100.0%	All prompt displays occurred in the intervention group.
Level 2 pause prompts shown	131	22.7%	Shown during Warning-level sessions.
Level 3 breathing prompts shown	445	77.3%	Most shown prompts were tied to Critical sessions.
Continue responses	242	42.0%	Most common response after a shown prompt.
Take-break responses	170	29.5%	Nearly one-third of prompt responses indicated a break.
Dismissed responses	112	19.4%	A notable minority dismissed prompts.
View-dashboard responses	52	9.0%	Dashboard viewing was the least common response.

Prompt responses suggest that the intervention was engaged with in the simulated logs, but users often continued rather than taking a break. This pattern supports recommendations for prompt timing, prompt burden management, and clearer dashboard value.

Table 16. Reliability Event Summary

Reliability Event Type	Frequency	Interpretation
SERVICE_STARTED	2,051	A service-start event was paired with each exported session.
TARGET_FOREGROUND	2,051	Each session had a target-app foreground event.
TARGET_BACKGROUND	2,051	Each session had a target-app background event.
SESSION_FINALIZED	2,051	Each session was finalized and saved.
HIGH_OOV	345	High out-of-vocabulary or unresolved text caused sentiment-unreliable fallback.

The reliability log is internally consistent because each exported session is matched by a service-start, foreground, background, and finalization event. The recurring reliability issue is HIGH_OOV, matching the 345 sentiment-unreliable sessions.

4.1.7. Research Question 3: Baseline Convergent Association Between DSI and Self-Reported Doomscrolling

The third research question asks what baseline convergent association exists between the fixed-prior Week 1 DSI and participants' self-reported doomscrolling scores among participants with at least three sentiment-reliable Week 1 sessions. The export supports both DSI coverage analysis and the baseline self-report correlation because Week 1 Doomscrolling Scale scores are present for all participants.

Table 17. DSI Coverage and Sentiment-Reliable Session Summary

Coverage Indicator	Intervention Group	Control Group	Total	Interpretation Notes
Participants with ≥ 3 sentiment-reliable Week 1 sessions	25	25	50	All participants met the log coverage rule.
Week 1 sentiment-reliable sessions	427	434	861	Enough retained sessions for week-level log aggregation.
Week 1 sentiment-unreliable sessions	94	82	176	Excluded from NSD and DSI calculations requiring sentiment.
Undefined Week 1 DSI cases	0	0	0	No participant failed the log-coverage rule.

The DSI coverage result is favorable for the simulated logs: all participants had at least three reliable Week 1 sessions, and the Week 1 reliable-session count was high enough to compute log-based baseline aggregates.

Table 18. Baseline Component and Composite Correlation with Doomscrolling Scale

Baseline Variable	N_effective	Correlation Type	Coefficient	p-value	Interpretation
Mean daily session duration	50	Pearson r	0.966	p less than 0.001	Strong positive association.
Mean daily video dwell time	50	Pearson r	-0.916	p less than 0.001	Strong negative association in this scoring direction.
Mean daily NSD	50	Pearson r	0.980	p less than 0.001	Strong positive association.
Composite Week 1 DSI	50	Pearson r	0.989	p less than 0.001	Strongest baseline association.

The composite Week 1 DSI showed the strongest association with the baseline Doomscrolling Scale. This supports preliminary estimator plausibility in the simulated dataset because the composite score related more strongly to self-reported doomscrolling than

session duration alone. The finding should still be framed as convergent association, not diagnostic validation, because the system estimates risk from observable proxies and the study remains pilot-scale.

4.1.8. ISO/IEC 25010 Software-Quality Criteria and Per-Category User Results

The fourth research question asks how users evaluate the system using ISO/IEC 25010 and TAM. The export contains 50 ISO/IEC 25010 survey rows, 50 SUS response rows, and 50 TAM response rows.

Table 19. ISO/IEC 25010 User Evaluation Summary

Quality Characteristic	Items	Mean	SD	Target	Interpretation
Functional Suitability	3	3.97	0.35	≥ 3.50	Met; favorable.
Performance Efficiency	3	3.83	0.38	≥ 3.50	Met; favorable.
Reliability	3	3.91	0.40	≥ 3.50	Met; favorable.

Table 20. System Usability Scale Result

Usability Measure	Analyzed n	Mean Score	Target	Interpretation
System Usability Scale (0-100)	50	76.35	≥ 70	Met; favorable usability.

Table 21. Technology Acceptance Model User Evaluation Summary

TAM Construct	Items	Mean	SD	Target	Interpretation
Perceived Usefulness	6	4.01	0.40	≥ 3.50	Met; favorable.
Perceived Ease of Use	6	4.06	0.31	≥ 3.50	Met; favorable.

All user-evaluation means met the study-defined favorable thresholds. Users rated the system favorably on functional suitability, performance efficiency, reliability, usability, perceived usefulness, and perceived ease of use. The low internal-consistency coefficients reported below indicate that these favorable means should still be interpreted cautiously at the subscale level.

4.1.9. Reliability Analysis

Table 22. Cronbach's Alpha Results for User Survey Subscales

Subscale	Items	Cronbach's Alpha	Interpretation
Functional Suitability	3	0.106	Below 0.70; interpret subscale cautiously.
Performance Efficiency	3	0.050	Below 0.70; interpret subscale cautiously.
Reliability	3	0.172	Below 0.70; interpret subscale cautiously.
System Usability Scale	10	0.207	Below 0.70 after SUS item scoring; interpret cautiously.
Perceived Usefulness	6	0.687	Slightly below 0.70; near-threshold coherence.
Perceived Ease of Use	6	0.254	Below 0.70; interpret subscale cautiously.

Cronbach's alpha results were mostly below the conventional 0.70 benchmark. This weak internal consistency may reflect the small pilot sample, heterogeneous item wording, and limited number of items in the researcher-developed ISO/IEC 25010 subscales. The Perceived Usefulness subscale was closest to the threshold at 0.687. The favorable user means should therefore be reported alongside this reliability caution.

4.1.10. Algorithm Complexity Analysis

Table 23. Algorithm Complexity Analysis

Algorithm / Procedure	Input Basis	Complexity Reported	Interpretation Notes
Accessibility event processing	Event stream and current UI tree size.	$O(e + n)$ per event batch	Processing grows with captured events and traversed nodes.
Sentiment and fallback routing	Extracted text tokens or no-text/unreliable marker.	$O(t)$ for text path	Runtime grows with token count; fallback drops sentiment input.
Fuzzy risk estimation	Three inputs and 27 fixed rules, or two inputs and fallback rules.	$O(1)$ for fixed rule base	Rule count is constant in the fielded design.
Local aggregation and export	Stored session, prompt, reliability, and personalization rows.	$O(r)$ over exported rows	Export cost grows linearly with retained aggregate records.

The complexity profile is appropriate for an on-device pilot system because the fuzzy inference rule base is fixed and the exported data are aggregate records. The main

scalability risks are not the rule engine itself but platform extraction stability, text length, and device-level background service behavior.

4.1.11. Research Question 5: Subject Matter Expert Evaluation

The fifth research question asks how subject matter experts evaluate the system's technical design, privacy safeguards, heuristic logic, and intervention structure. The SME files are present in the export, but the rating and comment fields are blank; therefore, SME evaluation remains pending.

Table 24. Subject Matter Expert Evaluation Summary

Evaluation Area	Experts	Mean	SD	Target	Interpretation
Technical soundness of risk-estimation framework	3 pending	Pending	Pending	≥ 4.00	Awaiting completed ratings.
Plausibility of input ranges	3 pending	Pending	Pending	≥ 4.00	Awaiting completed ratings.
Internal coherence of fuzzy rule base	3 pending	Pending	Pending	≥ 4.00	Awaiting completed ratings.
Privacy and architecture quality	3 pending	Pending	Pending	≥ 4.00	Awaiting completed ratings.
Intervention design appropriateness	3 pending	Pending	Pending	≥ 4.00	Awaiting completed ratings.
Overall expert evaluation	3 pending	Pending	Pending	≥ 4.00	Awaiting completed ratings.

The SME objective remains open. The log and user data can show operational behavior, self-report association, and user acceptance, but expert appraisal is still needed to judge whether the estimator design, thresholds, privacy safeguards, and intervention structure are technically appropriate.

4.1.12. Open-Ended Feedback

Table 25. Open-Ended Feedback Coding Summary

Theme / Issue Category	Frequency	Representative Concern	Design Implication
Prompt timing and interruption	14	Prompts sometimes appeared while users were deciding whether to continue.	Add quiet hours, gentler frequency, and better timing controls.
Permission and onboarding clarity	12	Users wanted clearer permission and setup explanation.	Improve onboarding instructions and permission rationale.
Dashboard usefulness	12	Users wanted clearer safe, warning, and critical explanations.	Add plain-language labels and weekly trends.
Privacy and trust	6	Users were comfortable with study-code exports if raw captions and screenshots were not saved.	Keep aggregate-only export and communicate it clearly.
Technical reliability	6	Some sessions seemed harder for the app to interpret.	Expose sentiment-reliability status and platform-detection limitations.

The most frequent user feedback theme was prompt timing and interruption, followed by permission/onboarding clarity and dashboard usefulness. These themes are consistent with the prompt-response logs, where continuing was the most common response and dashboard viewing was the least common response.

4.1.13. Overall Evaluation Results

Table 26. Overall Results Summary

Evaluation Area	Evidence Source	Overall Result	Interpretation Status
Implementation and testing	Architecture matrix, exported sessions, reliability events, and prompt events	Logging, scoring, fallback flagging, and prompt-event generation are visible in the export.	Log-supported
Behavioral and self-report outcomes	Week 1 to Week 2 logs and Doomscrolling Scale scores	No reliable between-group difference was observed in log metrics, but self-reported doomscrolling decreased more in the intervention group.	Mixed but partly favorable
Estimator plausibility	Week 1 DSI coverage and Doomscrolling Scale correlation	All participants met the log coverage rule, and composite DSI correlated	Supported as preliminary plausibility

Evaluation Area	Evidence Source	Overall Result	Interpretation Status
		strongly with baseline self-report.	
User evaluation	ISO/IEC 25010, SUS, TAM, and open-ended feedback	All user mean targets were met; reliability coefficients were mostly below 0.70 and require caution.	Supported with reliability caution
Expert evaluation	SME rating form and SME comments	SME files are present but blank.	Pending

4.2. Narrative Interpretation of Results

The simulated export shows that REDU can produce a privacy-preserving aggregate log dataset across the three target short-form platforms. The strongest evidence is operational: all 50 participants have Week 1 and Week 2 logs, all sessions have risk scores and risk levels, all sessions have reliability-event traces, and prompt events are limited to the intervention group as intended.

The behavioral comparison does not support a strong short-term difference between the intervention and control groups in the available simulated logs. The intervention group showed a small increase in mean daily raw elapsed session duration, while the control group showed a small decrease, but the adjusted result was not statistically reliable. Dwell time and NSD showed negligible between-group differences. The self-reported Doomscrolling Scale, however, decreased more in the intervention group than in the control group. This creates a mixed finding: the intervention was associated with a favorable self-report change, but not with a clear short-term reduction in logged use metrics.

The DSI coverage result is favorable because every participant met the minimum Week 1 sentiment-reliable session rule. The composite Week 1 DSI also showed a strong positive correlation with baseline self-reported doomscrolling. This supports preliminary estimator plausibility, but it should not be overstated as diagnostic validation.

User evaluation results were favorable at the mean-score level. ISO/IEC 25010, SUS, and TAM results met the study-defined targets, and the open-ended feedback showed that users generally recognized the value of break reminders, privacy safeguards, and activity summaries. The main weaknesses were prompt timing, onboarding clarity, dashboard explanation, and technical reliability visibility. SME evaluation remains pending.

4.3. Summary of Findings

Table 27. Summary of Findings by Research Question

RQ	Focus	Key Result	Interpretation Status
1	Privacy-preserving architecture and estimation framework	2,051 aggregate sessions, 1,706 reliable sessions, 345 fallback-flagged sessions, and 576 intervention-only prompt displays were exported without raw content.	Log-supported objective
2	Week 1 to Week 2 behavioral and self-report changes	No adjusted between-group difference was observed for log metrics, but the Doomscrolling Scale improved more in the intervention group.	Mixed; self-report favorable
3	Baseline DSI and Doomscrolling Scale association	All 50 participants met Week 1 reliable-session coverage; composite DSI correlated with DSS at $r=0.989$, p less than 0.001.	Preliminary estimator plausibility
4	User evaluation through ISO/IEC 25010 and TAM	ISO, SUS, and TAM means met targets; open-ended feedback identified prompt timing and onboarding as top issues.	Supported with reliability caution
5	Subject matter expert evaluation	SME files are present but blank.	Pending

Overall, the simulated export supports REDU’s aggregate logging and rule-based risk-estimation workflow, provides preliminary self-report convergence evidence, and

shows favorable user acceptance at the mean-score level. It does not yet provide completed SME appraisal, and the behavioral-log results remain mixed rather than clearly favorable.

Chapter 5

CONCLUSION

This chapter presents the summary of findings and conclusions of the study based on the simulated export analyzed in Chapter 4. The conclusions are bounded by the available data. The export supports conclusions about aggregate logging, fallback flagging, prompt-event generation, pre/post self-reported doomscrolling, baseline convergent association, ISO/IEC 25010 user ratings, SUS usability, TAM acceptance, and coded user feedback. SME ratings and SME comments remain pending because the expert files are present but blank.

5.1. Summary of Findings

The study was guided by five research questions. The first question focused on the design and implementation of a privacy-preserving architecture and estimation framework. The second question examined short-term Week 1-to-Week 2 changes in logged usage metrics and self-reported doomscrolling scores between the intervention group and the logging-only control group. The third question examined the baseline convergent association between the fixed-prior Week 1 DSI and participants' self-reported doomscrolling scores. The fourth question evaluated the system from the users' perspective using ISO/IEC 25010 and TAM. The fifth question obtained subject matter expert evaluation of the technical design, privacy safeguards, heuristic logic, and intervention structure.

Table 28. Summary of Findings

RQ	Research Question Focus	Finding from Chapter 4	Conclusion Direction
1	Architecture and estimation framework	The export contained 2,051 aggregate sessions, 1,706	Objective supported by logs

RQ	Research Question Focus	Finding from Chapter 4	Conclusion Direction
		sentiment-reliable sessions, 345 fallback cases, and 576 intervention-only prompt displays.	
2	Week 1 to Week 2 behavioral and self-report changes	Log metrics showed no reliable between-group reduction, but the Doomscrolling Scale decreased more in the intervention group than in the control group.	Mixed; self-report favorable
3	Baseline DSI and Doomscrolling Scale association	All 50 participants met the Week 1 reliable-session coverage rule, and composite DSI correlated strongly with baseline DSS.	Preliminary estimator plausibility
4	User ISO/IEC 25010 and TAM evaluation	ISO/IEC 25010, SUS, and TAM means met targets, while alpha coefficients were mostly below 0.70.	User acceptance favorable with reliability caution
5	SME evaluation	SME files were present but rating and comment fields were blank.	Expert appraisal pending

The strongest findings are that REDU produced a complete aggregate logging workflow, met the user-evaluation mean targets, and showed a strong baseline association between composite DSI and self-reported doomscrolling. The short-term behavioral logs did not show reliable between-group improvement, so the self-report improvement should be interpreted cautiously rather than as definitive behavioral effectiveness.

5.2. Conclusions Based on Research Objectives and Questions

For the first research question, the simulated export supports the conclusion that the implemented architecture can produce privacy-preserving aggregate records for the target study workflow. The exported files contain study codes, platform labels, timing metrics, risk scores, reliability codes, and prompt actions without raw captions, comments, screenshots, or directly identifying content. The system also produced fallback flags for 345 sentiment-

unreliable sessions and preserved analyzable NSD values for 1,706 sessions. Therefore, the architecture and estimation framework are supported at the exported-log level, although SME review is still needed for expert appraisal of the design.

For the second research question, the available log metrics do not show a reliable short-term between-group behavioral difference. The intervention group's mean daily raw elapsed session duration increased by 0.59 minutes while the control group's decreased by 0.38 minutes, but the Holm-adjusted p -value was 0.200 and the effect-size confidence interval crossed zero. Dwell time and NSD differences were smaller. However, the Doomscrolling Scale mean score decreased more in the intervention group than in the control group, with an adjusted p -value of 0.005. Thus, the simulated data support a favorable self-report change but not a corresponding reliable reduction in logged behavior within the two-week pilot window.

For the third research question, the coverage requirement for DSI computation was met because all 50 participants had at least three sentiment-reliable Week 1 sessions. The composite Week 1 DSI correlated strongly with baseline Doomscrolling Scale scores ($r = 0.989$, $p < 0.001$), and it showed a stronger association than session duration, dwell time, or NSD alone. This supports preliminary estimator plausibility in the simulated dataset, but it remains a convergent association finding rather than diagnostic validation.

For the fourth research question, users evaluated the system favorably at the mean-score level. Functional Suitability (3.97), Performance Efficiency (3.83), Reliability (3.91), SUS usability (76.35), Perceived Usefulness (4.01), and Perceived Ease of Use (4.06) all met the study-defined targets. Open-ended feedback identified prompt timing, permission/onboarding clarity, and dashboard usefulness as the most frequent user concerns. Because most Cronbach's alpha values were below 0.70, the favorable means should be interpreted with caution and supported by future instrument refinement.

For the fifth research question, no conclusion can be made about subject matter expert evaluation from the current export. Expert ratings and comments are still required to judge the technical soundness of the risk-estimation framework, input ranges, fuzzy rule base, privacy safeguards, and intervention structure.

The overall conclusion is that REDU is supported as a privacy-preserving, aggregate-log-producing prototype with visible risk scoring, fallback handling, prompt-event recording, preliminary self-report convergence, and favorable user-evaluation means. It should not yet be described as behaviorally effective in logged use or expert-approved. The study remains a pilot-scale software evaluation rather than a clinical or diagnostic validation study.

Chapter 6

RECOMMENDATIONS

This chapter presents recommendations for system improvement, future research, and future implementation or deployment. The recommendations are based on the simulated export analyzed in Chapter 4. The available user data support recommendations about prompt timing, onboarding clarity, dashboard usefulness, privacy communication, survey reliability, and sentiment-reliability coverage. SME recommendations remain pending until the blank expert rating and comment files are completed.

6.1. Recommendations for System Improvement

System-improvement recommendations are tied to verified implementation results, user survey results, and coded user feedback.

Table 29. System Improvement Recommendations

Improvement Area	Condition Observed in Chapter 4	Recommendation
Platform extraction robustness	All platforms produced logs, but sentiment reliability varied: TikTok 82.6%, Facebook 86.3%, and Instagram 81.9%.	Maintain platform-specific extraction tests and update selectors when target-app UI structures change.
Android compatibility and performance	Performance Efficiency met the target with mean 3.83, but technical reliability appeared in 6 user feedback rows.	Add CPU, memory, battery, and background-service survival logs, and expose reliability status in the dashboard.
Sentiment and fallback handling	345 of 2,051 sessions, or 16.8%, were sentiment-unreliable and marked HIGH_OOV.	Expand the Filipino/Taglish lexicon review process and improve reliability diagnostics for high-OOV sessions.
Prompt timing and burden management	Prompt timing and interruption was the most frequent feedback theme with 14 responses.	Add quiet hours, gentler prompt-frequency controls, and clearer cooldown behavior.

Improvement Area	Condition Observed in Chapter 4	Recommendation
Dashboard and user-facing transparency	Dashboard usefulness appeared in 12 feedback rows, and only 9.0% of prompt responses selected view-dashboard.	Add weekly trends, plain-language Safe/Warning/Critical explanations, and clearer DSI component summaries.

The highest-priority system improvements are prompt timing and dashboard clarity because these were the most frequent user-facing concerns. Sentiment-reliability handling remains technically important because unresolved or high-OOV sessions directly reduce NSD and DSI interpretability.

6.2. Recommendations for Future Development

Future-development recommendations should address the prototype limitations shown by the logs, user survey results, and coded feedback.

Table 30. Future Development Recommendations

Development Area	Observed Basis	Recommendation
Risk-estimation model refinement	DSI coverage was complete and DSI correlated strongly with baseline DSS, but 16.8% of sessions lacked reliable sentiment.	Retain the interpretable fuzzy model while adding better missing-sentiment diagnostics and sensitivity summaries.
Platform support expansion	All three target platforms were represented, but extraction success differed by platform.	Build automated regression checks for each target platform and document platform-specific failure modes.
User interface and dashboard iteration	SUS met the target at 76.35, but dashboard usefulness and prompt timing were recurring feedback themes.	Iterate the dashboard and prompt flow around the most frequent feedback categories.
Data export and reporting workflow	User files are now present, but SME rating and comment files remain blank.	Keep the expanded export format and complete the SME data collection fields before final defense.

Future development should preserve the expanded export package because it now supports Chapter 4 user analyses. The remaining export gap is expert review completion.

6.3. Recommendations for Future Researchers

Future-research recommendations are grounded in the final study limitations, sample coverage, statistical results, and survey reliability findings.

Table 31. Future Research Recommendations

Research Area	Observed Basis	Recommendation
Sample size and group balance	The simulated export had balanced groups of 25 participants each and complete week-level log and user-response coverage.	Use at least the same balanced structure in field deployment, but plan for attrition and device incompatibility.
Deployment duration	Log metrics did not show reliable Week 1-to-Week 2 behavioral reduction, while DSS improved.	Use a longer deployment window if the goal is to evaluate sustained behavior change in logs.
Validation anchors	Composite Week 1 DSI correlated strongly with baseline DSS at $r=0.989$.	Replicate the convergent association with real field data and test whether DSI predicts later self-report or intervention response.
Survey reliability	Most alpha coefficients were below 0.70; Perceived Usefulness was closest at 0.687.	Refine item wording, add or revise weak items, and pilot the instrument before final deployment.
Qualitative feedback	The top themes were prompt timing (14), onboarding clarity (12), and dashboard usefulness (12).	Use follow-up interviews or structured probes to explain why these issues recur.

Future researchers should treat the current findings as pilot evidence. The most important next step is to verify whether the favorable self-report and user-evaluation results remain stable in field data and whether they eventually align with behavioral log changes.

6.4. Recommendations for Future Implementation and Deployment

Future implementation and deployment recommendations are based on verified deployment constraints, privacy considerations, prompt interaction behavior, and user feedback.

Table 32. Future Implementation and Deployment Recommendations

Deployment Area	Observed Basis	Recommendation
Privacy and data governance	Privacy and trust feedback was favorable but still appeared as a distinct theme in 6 responses.	Keep the aggregate-only export design and explain clearly that raw captions and screenshots are not saved.

Deployment Area	Observed Basis	Recommendation
Participant onboarding	Permission and onboarding clarity appeared in 12 feedback rows.	Add clearer setup screens, permission rationale, and a post-setup checklist.
Monitoring and support	Every session had service-start, foreground, background, and finalization events; 345 HIGH_OOV events were recorded.	Monitor reliability-event counts during deployment and contact participants when service or extraction issues recur.
Institutional or classroom deployment	User acceptance means were favorable, but SME acceptability is still pending.	Avoid classroom-scale rollout until SME ratings and ethics-facing deployment safeguards are completed.
Maintenance and platform updates	Platform extraction success differed across target apps.	Schedule platform-maintenance checks before and during field deployment to catch app UI changes early.

The system may be responsibly improved as a privacy-preserving prototype. Broader deployment should wait until SME review is complete and until prompt timing, onboarding clarity, dashboard explanation, and survey reliability concerns are addressed.

REFERENCES

- Ababat, E. A., Alcantara, G. J., Jumamoy, R. T., Nuñez, R. M. M., & Suba, J. G. (2024,). *Excessive usage of internet on selected senior high school students at Bestlink College of the Philippines*. <https://ojs.aaresearchindex.com/index.php/aasgbcpjmra/article/view/13281>
- Adnan, A., Irvine, R. E., Williams, A., Harris, M., & Antonacci, G. (2025). Improving acceptability of mHealth apps—The use of the technology acceptance model to assess the acceptability of mHealth apps: Systematic review. *Journal of Medical Internet Research*, 27, e66432. <https://doi.org/10.2196/66432>
- Aggarwal, D., Sharma, D., & Saxena, A. B. (2024). Developing digital mindfulness for a thoughtful and sensible technology usage for sustainability. In *Research Updates in Mathematics and Computer Science Vol. 6* (pp. 1–10). B P International. <https://doi.org/10.9734/bpi/rumcs/v6/12466F>
- Ahmed, M. S., & Ahmed, N. (2023). A Fast and Minimal System to Identify Depression Using Smartphones: Explainable Machine Learning–Based Approach. *JMIR Formative Research*, 7, e28848. <https://doi.org/10.2196/28848>
- Antezana, G., Venning, A., Smith, D., & Bidargaddi, N. (2022). Understanding what we know so far about young people's engagement with wellbeing apps. A scoping review and narrative synthesis. *DIGITAL HEALTH*, 8, 205520762211441. <https://doi.org/10.1177/20552076221144104>
- Apple Inc. (2025,). *Get started with Screen Time on iPhone*. <https://support.apple.com/en-me/guide/iphone/iphbfa595995/ios>

- Augusma, A., Vaufreydaz, D., & Letué, F. (2023). Multimodal Group Emotion Recognition In-the-wild Using Privacy-Compliant Features. *INTERNATIONAL CONFERENCE ON MULTIMODAL INTERACTION*, 750–754. <https://doi.org/10.1145/3577190.3616546>
- Azam, M. H., Hasan, M. H., Hassan, S., & Abdulkadir, S. J. (2021). A novel approach to generate type-1 fuzzy triangular and trapezoidal membership functions to improve the classification accuracy. *Symmetry*, 13(10), 1932. <https://doi.org/10.3390/sym13101932>
- Bautista, R. A. L., Demanawa, M. L. S., Pedrigosa, J. L. D., & Culajara, C. L. B. (2024). Boredom proneness and social media usage of college students. *American Journal of Human Psychology*, 2(1), 121–131. <https://doi.org/10.54536/ajhp.v2i1.2825>
- Buchanan, K., Aknin, L. B., Lotun, S., & Sandstrom, G. M. (2021). Brief exposure to social media during the COVID-19 pandemic: Doom-scrolling has negative emotional consequences, but kindness-scrolling does not. *PLOS ONE*, 16(10), e257728. <https://doi.org/10.1371/journal.pone.0257728>
- Buekers, J., Chernova, J., Koch, S., Marchena, J., Lemos, J., Becker, C., Bonci, T., Braun, J., Brittain, G., Buckley, E., Caulfield, B., Del Din, S., Demeyer, H., Frei, A., Gazit, E., Hausdorff, J. M., Ionescu, A., Jansen, C.-P., Klenk, J., ... Garcia-Aymerich, J. (2025). Digital assessment of real-world walking in people with impaired mobility: How many hours and days are needed?. *International Journal of Behavioral Nutrition and Physical Activity*, 22(1). <https://doi.org/10.1186/s12966-025-01851-3>
- Canila, J. F., Cruz, M. C., Fernandez, A. T., Pasao, G. M., & Regalado, M. M. (2023). *Perceived effects of Tiktok consumption on the attention span of DLSMHSI senior high school students S.Y. 2022–2023*. https://greenprints.dlshsi.edu.ph/grade_12/517/
- Cardoso, C. J., Acosta, N. J., Borja, M. S. T., Ichikawa, T. J. G., Nayos, M. A. R., Ocampo, B. J. A., & Cayanan, E. C. (2024). Impact of internet usage in the declining attention span of Radiologic Technology students: An assessment. *International Journal of Education*

- Humanities and Social Science*, 7(3), 166–173. <https://doi.org/10.54922/IJEHSS.2024.0704>
- Casalino, G., Castellano, G., Castiello, C., & Mencar, C. (2022). Effect of fuzziness in fuzzy rule-based classifiers defined by strong fuzzy partitions and winner-takes-all inference. *Soft Computing*, 26(14), 6519–6527. <https://doi.org/10.1007/s00500-022-07128-2>
- Castillo, J. M. C., Garcia, L. L., Abalos, E., & Locsin, R. C. (2022). Living alone and using social media technologies: The experience of Filipino older adults during the COVID-19 pandemic. *Nursing Inquiry*, 29(3). <https://doi.org/10.1111/nin.12460>
- Chen, Y., Fu, Y., Chen, Z., Radesky, J., & Hiniker, A. (2024). The engagement-prolonging designs teens encounter on very large online platforms. *Arxiv Preprint Arxiv:2411.12083*. <https://arxiv.org/abs/2411.12083>
- Chen, Y.-C., Lee, Y.-J., Kao, K.-C., Tsai, J., Liang, E.-C., Chiu, W.-C., Shih, F., & Chang, Y.-J. (2023). Are You Killing Time? Predicting Smartphone Users' Time-killing Moments via Fusion of Smartphone Sensor Data and Screenshots. *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*, 1–19. <https://doi.org/10.1145/3544548.3580689>
- Cho, H., Choi, D., Kim, D., Kang, W. J., Choe, E. K., & Lee, S.-J. (2021). Reflect, not regret: Understanding regretful smartphone use with app feature-level analysis. *Proceedings of the ACM on Human-Computer Interaction*, 5(CSCW2), 1–36. <https://doi.org/10.1145/3479600>
- Cleofas, J. V., Dayrit, J. C. S., & Albao, B. T. (2022). Problematic versus reflective use: Types of social media use as determinants of mental health among young Filipino undergraduates. *Health Promotion Perspectives*, 12(1), 85–91. <https://doi.org/10.34172/hpp.2022.11>
- Co, N. A., Estuar, M. R. J., Tan, H. C., Tan, A. S., Abao, R., & Aureus, J. (2022). Development of bilingual sentiment and emotion text classification models from COVID-19

- vaccination tweets in the Philippines. In *Social Computing and Social Media: Design, User Experience and Impact* (pp. 247–266). Springer International Publishing. https://doi.org/10.1007/978-3-031-05061-9_18
- Cortiñas-Lorenzo, K., & Lacey, G. (2024). Toward explainable affective computing: A review. *IEEE Transactions on Neural Networks and Learning Systems*, 35(10), 13101–13121. <https://doi.org/10.1109/TNNLS.2023.3270027>
- Creswell, J. W., & Creswell, J. D. (2022). *Research design: Qualitative, quantitative, and mixed methods approaches* (6th ed.). SAGE Publications.
- Cruz, C. A., & Balahadia, F. (2022). Analyzing public concern responses for formulating ordinances and laws using sentiment analysis through VADER application. *International Journal of Computing Sciences Research*, 6, 842–856. <https://doi.org/10.25147/ijcsr.2017.001.1.77>
- Das, R., & Singh, T. D. (2023). Multimodal sentiment analysis: A survey of methods, trends, and challenges. *ACM Computing Surveys*, 55(13s), 1–38. <https://doi.org/10.1145/3586075>
- Dogan, O. (2021). Process mining technology selection with spherical fuzzy AHP and sensitivity analysis. *Expert Systems with Applications*, 178, 114999. <https://doi.org/10.1016/j.eswa.2021.114999>
- Fiedler, J., Bergmann, M. R., Sell, S., Woll, A., & Stetter, B. J. (2024). Just-in-Time Adaptive Interventions for Behavior Change in Physiological Health Outcomes and the Use Case for Knee Osteoarthritis: Systematic Review. *Journal of Medical Internet Research*, 26, e54119. <https://doi.org/10.2196/54119>
- Gagalang, J. L. (2021). Exploring social media use of Filipino learners: How it impacts reading attitudes and competence. *Linguistics and Culture Review*, 6, 275–290. <https://doi.org/10.21744/lingcure.v6ns2.2034>

- Genugten, C. R. van, Thong, M. S. Y., Ballegooijen, W. van, Kleiboer, A. M., Spruijt-Metz, D., Smit, A. C., Sprangers, M. A. G., Terhorst, Y., & Riper, H. (2025). Beyond the current state of just-in-time adaptive interventions in mental health: A qualitative systematic review. *Frontiers in Digital Health*, 7. <https://doi.org/10.3389/fdgth.2025.1460167>
- Google. (2024, May). *Manage how you spend time on your Android phone with Digital Wellbeing*. <https://support.google.com/android/answer/9346420>
- Hashmi, E., Yayilgan, S. Y., & Shaikh, S. (2024). Augmenting sentiment prediction capabilities for code-mixed tweets with multilingual transformers. *Social Network Analysis and Mining*, 14(1). <https://doi.org/10.1007/s13278-024-01245-6>
- Hawwa, B., Arif, Z., & Taj, S. (2025). Impact of Social Media Addiction on Anxiety: Mediating Role of Doom-scrolling Among Young Adults. *Social Sciences & Humanity Research Review*, 3(3). <https://doi.org/10.63468/sshrr.084>
- Hsu, T. C., Whelan, P., Gandrup, J., Armitage, C. J., Cordingley, L., & McBeth, J. (2025). Personalized interventions for behaviour change: A scoping review of just-in-time adaptive interventions. *British Journal of Health Psychology*, 30(1). <https://doi.org/10.1111/bjhp.12766>
- Hussain, M., Chen, C., Anwar, M., Ghorashi, S. A., Ahmed, A., Malik, M. S. A., & Yamin, I. (2025). Adaptive multitask emotion recognition and sentiment analysis using resource-constrained MobileBERT and DistilBERT: An efficient approach for edge devices. *Peerj Computer Science*, 11, e3287. <https://doi.org/10.7717/peerj-cs.3287>
- Hutto, C. J., & Gilbert, E. (2014). VADER: A Parsimonious Rule-Based Model for Sentiment Analysis of Social Media Text. *Proceedings of the International AAAI Conference on Web and Social Media*, 8(1), 216–225. <https://doi.org/10.1609/icwsm.v8i1.14550>

- Hyzy, M., Bond, R., Mulvenna, M., Bai, L., Dix, A., Leigh, S., & Hunt, S. (2022). System Usability Scale benchmarking for digital health apps: Meta-analysis. *JMIR Mhealth and Uhealth*, 10(8), e37290. <https://doi.org/10.2196/37290>
- Ikegaya, M., Foo, J. C., Murata, T., Oshima, K., & Kim, J. (2025). Using personalized intervention criteria in a mobile just-in-time adaptive intervention for increasing physical activity in university students: Pilot study. *JMIR Human Factors*, 12, e66750–e66750. <https://doi.org/10.2196/66750>
- International Organization for Standardization. (2023). *ISO/IEC 25010:2023 - Systems and software engineering — Systems and software Quality Requirements and Evaluation (SQuaRE) — Product quality model* (2nd ed.). International Organization for Standardization. <https://www.iso.org/standard/78176.html>
- Ismail, T., & Al Thani, D. (2022). Design and Evaluation of a Just-in-Time Adaptive Intervention (JITAI) to Reduce Sedentary Behavior at Work: Experimental Study. *JMIR Formative Research*, 6(1), e34309. <https://doi.org/10.2196/34309>
- Johnson, D. S., Hakobyan, O., Paletschek, J., & Drimalla, H. (2025). Explainable AI for audio and visual affective computing: A scoping review. *IEEE Transactions on Affective Computing*, 16(2), 518–536. <https://doi.org/10.1109/TAFFC.2024.3505269>
- Khairuddin, S. H., Hasan, M. H., Hashmani, M. A., & Azam, M. H. (2021). Generating Clustering-Based Interval Fuzzy Type-2 Triangular and Trapezoidal Membership Functions: A Structured Literature Review. *Symmetry*, 13(2), 239. <https://doi.org/10.3390/sym13020239>
- Khan, J., Ahmad, K., Jagatheesaperumal, S., & Sohn, K.-A. (2025). Textual variations in social media text processing applications: Challenges, solutions, and trends. *Artificial Intelligence Review*, 58(3). <https://doi.org/10.1007/s10462-024-11071-z>
- Krušić, L. (2024). Constructing a sentiment-annotated corpus of Austrian historical newspapers: Challenges, tools, and annotator experience. *Proceedings of the 4th Inter-*

- national Conference on Natural Language Processing for Digital Humanities*, 51–62.
<https://doi.org/10.18653/v1/2024.nlp4dh-1.6>
- Lanuza, J. M. H., Fallorina, R., & Cabbuag, S. (2021). *Understudied digital platforms in the Philippines*. https://internews.org/wp-content/uploads/2021/12/Internews_Understudied-Digital-Platforms-PH_December_2021.pdf
- Laudon, K. C., & Laudon, J. P. (2022). *Management information systems: Managing the digital firm* (17th ed.). Pearson.
- Lee, H., Park, J., & Lee, U. (2022). A systematic survey on Android API usage for data-driven analytics with smartphones. *ACM Computing Surveys*, 55(5), 1–38. <https://doi.org/10.1145/3530814>
- Lee, S. I., Koo, K., Lee, J. H., Lee, G., Jeong, S., O, S., & Kim, H. (2024). Vision transformer models for mobile/edge devices: a survey. *Multimedia Systems*, 30(2). <https://doi.org/10.1007/s00530-024-01312-0>
- Lim, S. G. S., Salazar, R. M. C., Tan, H. L. D., & Arcinas, M. (2025). Cognitive and social effects of a two-week digital detox on Filipino adolescents: A quasi-experimental study. *International Journal of Multidisciplinary: Applied Business and Education Research*, 6(5), 2353–2367. <https://doi.org/10.11594/ijmaber.06.05.21>
- Lokeshkumar, R., Mishra, O. A., & Kalra, S. (2021). Social media data analysis to predict mental state of users using machine learning techniques. *Journal of Education and Health Promotion*, 10(1). https://doi.org/10.4103/jehp.jehp_446_20
- Mair, J. L., Hayes, L. D., Campbell, A. K., Buchan, D. S., Easton, C., & Sculthorpe, N. (2022). A Personalized Smartphone-Delivered Just-in-time Adaptive Intervention (JitaBug) to Increase Physical Activity in Older Adults: Mixed Methods Feasibility Study. *JMIR Formative Research*, 6(4), e34662. <https://doi.org/10.2196/34662>
- Meinhardt, L.-M., Elhaidary, M., Colley, M., Rietzler, M., Rixen, J. O., Purohit, A. K., & Rukzio, E. (2025). Scrolling in the deep: Analysing contextual influences on inter-

- vention effectiveness during infinite scrolling on social media. *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*, 1–17. <https://doi.org/10.1145/3706598.3713187>
- Meyer, B. M., Depetrillo, P., Franco, J., Donahue, N., Fox, S. R., O’Leary, A., Loftness, B. C., Gurchiek, R. D., Buckley, M., Solomon, A. J., Ng, S. K., Cheney, N., Ceruolo, M., & McGinnis, R. S. (2022). How Much Data Is Enough? A Reliable Methodology to Examine Long-Term Wearable Data Acquisition in Gait and Postural Sway. *Sensors*, 22(18), 6982. <https://doi.org/10.3390/s22186982>
- Mike Jeizy P, P., Nimjames, F., Paul Ryan, A., Shan, V., & Bryan Louis G, L. (2024). Lost in the Feed: Exploring the Lived Experiences of Students on Doomscrolling. *International Journal of Advanced Multidisciplinary Research and Studies*, 4(3), 378–385. <https://doi.org/10.62225/2583049x.2024.4.3.2796>
- Mitsea, E., Drigas, A., & Skianis, C. (2023). Digitally Assisted Mindfulness in Training Self-Regulation Skills for Sustainable Mental Health: A Systematic Review. *Behavioral Sciences*, 13(12), 1008. <https://doi.org/10.3390/bs13121008>
- Mohammed, I., & Prasad, R. (2023). Building lexicon-based sentiment analysis model for low-resource languages. *Methodsx*, 11, 102460. <https://doi.org/10.1016/j.mex.2023.102460>
- Mosseri, A. (2021, December). *Raising the standard for protecting teens and supporting parents online*. <https://about.instagram.com/blog/spark/announcements/raising-the-standard-for-protecting-teens-and-supporting-parents-online>
- Muise, D., Markowitz, D. M., Reeves, B., Ram, N., & Robinson, T. N. (2024). (Mis)measurement of political content exposure within the smartphone ecosystem: Investigating common assumptions. *Journal of Quantitative Description: Digital Media*, 4. <https://doi.org/10.51685/jqd.2024.015>

- Nazir, M. K., Bilal, M., & Shongwe, S. C. (2026). Sentiment analysis for code-mixed low-resource languages: a systematic review of approaches, techniques, applications, challenges, and future directions. *Social Network Analysis and Mining*, 16(1). <https://doi.org/10.1007/s13278-026-01588-2>
- Pacol, C., & Palaoag, T. (2021). Bilingual lexicon approach to English-Filipino sentiment analysis of teaching performance. *IOP Conference Series: Materials Science and Engineering*, 1077(1), 12044. <https://doi.org/10.1088/1757-899X/1077/1/012044>
- Panaligan, A. (2021,). *The allure of doomscrolling and watching our fears onscreen*. <https://philstarlife.com/self/874584-allure-doomscrolling-watching-fears-onscreen>
- Perera, A., & Caldera, A. (2024). Sentiment Analysis of Code-Mixed Text: A Comprehensive Review. *JUCS - Journal of Universal Computer Science*, 30(2), 242–261. <https://doi.org/10.3897/jucs.98708>
- Pickering, L., Cohen, K., & De Baets, B. (2025). A Narrative Review on the Interpretability of Fuzzy Rule-Based Models from a Modern Interpretable Machine Learning Perspective. *International Journal of Fuzzy Systems*. <https://doi.org/10.1007/s40815-025-02022-z>
- Porebski, S. (2022). Evaluation of fuzzy membership functions for linguistic rule-based classifier focused on explainability, interpretability and reliability. *Expert Systems with Applications*, 199, 117116. <https://doi.org/10.1016/j.eswa.2022.117116>
- Qin, Y., Omar, B., & Musetti, A. (2022). The addiction behavior of short-form video app TikTok: The information quality and system quality perspective. *Frontiers in Psychology*, 13. <https://doi.org/10.3389/fpsyg.2022.932805>
- Qu, T., Tang, L., Peng, B., Yang, S., Yu, B., & Jia, J. (2025). Does your vision-language model get lost in the long video sampling dilemma?. *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 20889–20899. <https://openaccess>.

- thecvf.com/content/ICCV2025/html/Qu_Does_Your_Vision-Language_Model_Get_Lost_in_the_Long_Video_ICCV_2025_paper.html
- Rahmillah, F. I., Tariq, A., King, M., & Oviedo-Trespalacios, O. (2023). Evaluating the Effectiveness of Apps Designed to Reduce Mobile Phone Use and Prevent Maladaptive Mobile Phone Use: Multimethod Study. *Journal of Medical Internet Research*, 25, e42541. <https://doi.org/10.2196/42541>
- Rajeswari, N., & Vakkil, M. (2024). SCROLLING AND SUSTAINED ATTENTION: HOW SOCIAL MEDIA INFLUENCE SUSTAINED ATTENTION OF HIGH SCHOOL STUDENTS. *Shodhkosh: Journal of Visual and Performing Arts*, 5(6). <https://doi.org/10.29121/shodhkosh.v5.i6.2024.1726>
- Ratitch, B., Trigg, A., Majumder, M., Vlajnic, V., Rethemeier, N., & Nkulikiyinka, R. (2023). Clinical validation of novel digital measures: Statistical methods for reliability evaluation. *Digital Biomarkers*, 74–91. <https://doi.org/10.1159/000531054>
- Rixen, J. O., Meinhardt, L.-M., Glöckler, M., Ziegenbein, M.-L., Schlothauer, A., Colley, M., Rukzio, E., & Gugenheimer, J. (2023). The Loop and Reasons to Break It: Investigating Infinite Scrolling Behaviour in Social Media Applications and Reasons to Stop. *Proceedings of the ACM on Human-Computer Interaction*, 7(MHCI), 1–22. <https://doi.org/10.1145/3604275>
- Rodrigues, E. V. (2022). Doomscrolling – threat to Mental Health and Well-being: A Review. *International Journal of Nursing Research*, 8(4), 127–130. <https://doi.org/10.31690/ijnr.2022.v08i04.002>
- Roffarello, A. M., & De Russis, L. (2021). Understanding, Discovering, and Mitigating Habitual Smartphone Use in Young Adults. *ACM Transactions on Interactive Intelligent Systems*, 11(2), 1–34. <https://doi.org/10.1145/3447991>
- Satici, S. A., Gocet Tekin, E., Deniz, M. E., & Satici, B. (2023). Doomscrolling Scale: its Association with Personality Traits, Psychological Distress, Social Media Use, and

- Wellbeing. *Applied Research in Quality of Life*, 18(2), 833–847. <https://doi.org/10.1007/s11482-022-10110-7>
- Shahari, N. A., & Rasmani, K. (2024). Sensitivity Analysis for Fuzzy Similarity Measures. *ASM Science Journal*, 19, 1–10. <https://doi.org/10.32802/asmscj.2023.1180>
- Sharma, A., Sharma, K., & Kumar, A. (2023). Real-time emotional health detection using fine-tuned transfer networks with multimodal fusion. *Neural Computing and Applications*, 35(31), 22935–22948. <https://doi.org/10.1007/s00521-022-06913-2>
- Sharma, B., Lee, S. S., & Johnson, B. K. (2022). The dark at the end of the tunnel: Doomscrolling on social media newsfeeds. *Technology, Mind, And Behavior*, 3(1), 144–156. <https://doi.org/10.1037/tmb0000059>
- Sharshar, A., Khan, L. U., Ullah, W., & Guizani, M. (2025). Vision-Language Models for Edge Networks: A Comprehensive Survey. *IEEE Internet of Things Journal*, 12(16), 32701–32724. <https://doi.org/10.1109/jiot.2025.3579032>
- Shukla, A. K., & Muhuri, P. K. (2025). Deep belief network with fuzzy parameters and its membership function sensitivity analysis. *Neurocomputing*, 614, 128716. <https://doi.org/10.1016/j.neucom.2024.128716>
- Swathi, B., Veerabomma, S., Archana, M., Bhadru, D., Somu, N. L., & Bhavsingh, M. (2025). Edge-Centric IoT Health Monitoring: Optimizing Real-Time Responsiveness, Data Privacy, and Energy Efficiency. *2025 6th International Conference on Mobile Computing and Sustainable Informatics (ICMCSI)*, 354–361. <https://doi.org/10.1109/icmcsi64620.2025.10883456>
- Taskin, S., Kurtulus, H. Y., Satici, S. A., & Deniz, M. E. (2024). Doomscrolling and mental well-being in social media users: A serial mediation through mindfulness and secondary traumatic stress. *Journal of Community Psychology*, 52(3), 512–524. <https://doi.org/10.1002/jcop.23111>

- Teepe, G. W., Da Fonseca, A., Kleim, B., Jacobson, N. C., Salamanca Sanabria, A., Tudor Car, L., Fleisch, E., & Kowatsch, T. (2021). Just-in-Time Adaptive Mechanisms of Popular Mobile Apps for Individuals With Depression: Systematic App Search and Literature Review. *Journal of Medical Internet Research*, 23(9), e29412. <https://doi.org/10.2196/29412>
- Terzimehić, N., & Aragon-Hahner, S. (2022). I wish I had: Desired real-world activities instead of regretful smartphone use. *Proceedings of the 21st International Conference on Mobile and Ubiquitous Multimedia*, 47–52. <https://doi.org/10.1145/3568444.3568465>
- Tewari, A. (2023). mHealth systems need a privacy-by-design approach: Commentary on "Federated Machine Learning, Privacy-Enhancing Technologies, and Data Protection Laws in Medical Research: Scoping Review." *Journal of Medical Internet Research*, 25, e46700. *Journal of Medical Internet Research*, 25, e46700. <https://doi.org/10.2196/46700>
- Tho, C., Heryadi, Y., Kartowisastro, I. H., & Budiharto, W. (2021). A Comparison of Lexicon-based and Transformer-based Sentiment Analysis on Code-mixed of Low-Resource Languages. *2021 1st International Conference on Computer Science and Artificial Intelligence (ICCSAI)*, 81–85. <https://doi.org/10.1109/iccsai53272.2021.9609781>
- Tian, Y., Zhou, K., & Pelleg, D. (2021). What and How long: Prediction of Mobile App Engagement. *ACM Transactions on Information Systems*, 40(1), 1–38. <https://doi.org/10.1145/3464301>
- TikTok. (2024, June). *Well-being guide*. <https://www.tiktok.com/safety/en/well-being-guide>
- Vashishtha, S., Gupta, V., & Mittal, M. (2023). Sentiment analysis using fuzzy logic: A comprehensive literature review. *Wires Data Mining and Knowledge Discovery*, 13(5). <https://doi.org/10.1002/widm.1509>

- Vinogradova-Zinkevič, I. (2023). Comparative Sensitivity Analysis of Some Fuzzy AHP Methods. *Mathematics*, 11(24), 4984. <https://doi.org/10.3390/math11244984>
- Wang, J., Wu, Z., Choi, S. W., Sen, S., Yan, X., Miner, J. A., Sander, A. M., Lyden, A. K., Troost, J. P., & Carlozzi, N. E. (2023). The dosing of mobile-based just-in-time adaptive self-management prompts for caregivers: Preliminary findings from a pilot microrandomized study. *JMIR Formative Research*, 7, e43099. <https://doi.org/10.2196/43099>
- Wei, J., Yang, X., & Dong, Y. (2021). User-generated video emotion recognition based on key frames. *Multimedia Tools and Applications*, 80(9), 14343–14361. <https://doi.org/10.1007/s11042-020-10203-1>
- Wijayanti, R., & Arisal, A. (2021). Automatic Indonesian sentiment lexicon curation with sentiment valence tuning for social media sentiment analysis. *ACM Transactions on Asian and Low-Resource Language Information Processing*, 20(1), 1–16. <https://doi.org/10.1145/3425632>
- Yang, M.-J., Sutton, S. K., Hernandez, L. M., Jones, S. R., Wetter, D. W., Kumar, S., & Vinci, C. (2023). A Just-In-Time Adaptive intervention (JITAI) for smoking cessation: Feasibility and acceptability findings. *Addictive Behaviors*, 136, 107467. <https://doi.org/10.1016/j.addbeh.2022.107467>
- Yao, J., Tan, C. S., Lim, N., Tan, J., Chen, C., & Müller-Riemenschneider, F. (2021). Number of daily measurements needed to estimate habitual step count levels using wrist-worn trackers and smartphones in 212,048 adults. *Scientific Reports*, 11(1). <https://doi.org/10.1038/s41598-021-89141-3>
- Zamora, A. A., Regencia, Z. J. G., Crisostomo, M. E., Van Hal, G., & Baja, E. S. (2021). Effect of daily social media exposure on anxiety and depression disorders among cargo seafarers: a cross-sectional study. *International Maritime Health*, 72(1), 55–63. <https://doi.org/10.5603/imh.2021.0008>

- Zannettou, S., Nemes-Nemeth, O., Ayalon, O., Goetzen, A., Gummadi, K. P., Redmiles, E. M., & Roesner, F. (2024). Analyzing user engagement with TikTok's short format video recommendations using data donations. *Proceedings of the CHI Conference on Human Factors in Computing Systems*, 1–16. <https://doi.org/10.1145/3613904.3642433>
- Zenone, M., Ow, N., & Barbic, S. (2021). TikTok and public health: a proposed research agenda. *BMJ Global Health*, 6(11), e7648. <https://doi.org/10.1136/bmjgh-2021-007648>
- Zhang, H., & Xu, M. (2023). Recognition of emotions in user-generated videos through frame-level adaptation and emotion intensity learning. *IEEE Transactions on Multimedia*, 25, 881–891. <https://doi.org/10.1109/TMM.2021.3134167>
- Zhang, Y. (2023). Study on the influence of TikTok Use and individual socialization among adolescents. *Journal of Education, Humanities and Social Sciences*, 9, 204–210. <https://doi.org/10.54097/ehss.v9i.6454>
- Äyräväinen, L. E. M., Hinds, J., & Davidson, B. I. (2025). Disambiguating sentiment annotation: A mixed methods investigation of annotator experience and impact of instructions on annotator agreement. *PLOS One*, 20(12), e336269. <https://doi.org/10.1371/journal.pone.0336269>

APPENDICES

APPENDIX A

Complete 27-Rule Fuzzy Inference Base

The following table presents the complete 27-rule inference base used by the three-input fuzzy engine (Video Dwell Time, Negative Sentiment Density, Session Duration). The 27 rows arise from exhaustive coverage of the $3 \times 3 \times 3$ input-state combinations. Their Safe/Warning/Critical outputs are fixed study priors intended for transparent heuristic inference rather than empirically calibrated probabilities or clinical thresholds. The assignment logic is conservative: Safe is reserved for low-baseline or isolated single-signal states, Warning is the default for mixed or moderate states, and Critical is reserved for combinations where the main risk axes converge under the fixed study priors.

Table 33. Full 27-Rule Fuzzy Inference Base

No.	Video Dwell Time	Negative Sentiment Density	Session Duration	Risk Level
1	Low	Low	Low	Safe
2	Low	Low	Medium	Safe
3	Low	Low	High	Warning
4	Low	Medium	Low	Safe
5	Low	Medium	Medium	Warning
6	Low	Medium	High	Warning
7	Low	High	Low	Warning
8	Low	High	Medium	Warning
9	Low	High	High	Critical
10	Medium	Low	Low	Safe
11	Medium	Low	Medium	Warning
12	Medium	Low	High	Warning
13	Medium	Medium	Low	Warning
14	Medium	Medium	Medium	Warning
15	Medium	Medium	High	Critical
16	Medium	High	Low	Warning

Table 6. Full 27-Rule Fuzzy Inference Base (continued)

No.	Video Dwell Time	Negative Sentiment Density	Session Duration	Risk Level
17	Medium	High	Medium	Critical
18	Medium	High	High	Critical
19	High	Low	Low	Safe
20	High	Low	Medium	Warning
21	High	Low	High	Warning
22	High	Medium	Low	Warning
23	High	Medium	Medium	Critical
24	High	Medium	High	Critical
25	High	High	Low	Critical
26	High	High	Medium	Critical
27	High	High	High	Critical

Interpretive note: the table follows a mostly non-decreasing assignment pattern in Session Duration and Negative Sentiment Density, while using Video Dwell Time as a contextual modifier rather than a standalone high-risk trigger. One deliberate exception is retained as a precautionary prototype rule for rapid negative-content chaining. Low dwell combined with high Negative Sentiment Density and high session duration is classified as **Critical** because rapid movement does not reduce concern when the session is both prolonged and clearly negative-dominant overall. By contrast, Rule 19 (High Dwell, Low NSD, Low Session Duration) stays **Safe** because sustained attention to one item, without prolonged use or negative-dominant exposure, is not enough by itself to imply doomscrolling. This keeps dwell time as a modifier, not the main driver.

Operationally, the 27 consequents follow a compact assignment grammar rather than 27 unrelated judgments: no rule is **Critical** when both Session Duration and Negative Sentiment Density are Low; duration alone does not produce **Critical** when NSD remains Low; High NSD reaches **Critical** only when reinforced by another non-low axis; High Dwell Time alone is insufficient for **Critical**; and **Warning** remains the default mixed class when the signals are concerning but not yet strongly convergent.

Read in groups, the table has three layers. Rules 1, 2, 4, 10, and 19 remain **Safe** because at least two axes stay low. Rules 3, 5-8, 11-14, 16, and 20-22 remain **Warning** because they show mixed evidence without strong convergence between duration and negativity. Rules 9, 15, 17-18, and 23-27 become **Critical** because prolonged use and negative exposure rise together, or because High NSD is reinforced by High Dwell Time even before the session becomes long.

Within that Critical layer, the strongest challenge cases are still explainable. Rule 15 becomes **Critical** because a High-duration session with both other axes already at Medium no longer reflects time alone. Rule 23 becomes **Critical** because High Dwell Time, Medium NSD, and Medium Session Duration already show three non-low signals in the same direction. Rule 25 also remains **Critical** because a short session can still contain concentrated negative immersion when the viewed items are mostly negative and each item holds attention for longer.

APPENDIX B

Fallback Rules and Trigger Thresholds

Unless otherwise noted, the following thresholds are fixed study priors selected for reproducibility and conservative intervention behavior. They are design parameters, not empirically validated clinical or population cutoffs. Because the fallback engine retains two inputs with three linguistic levels each, exhaustive coverage requires $3^2 = 9$ rules.

In the current field-tested artifact, the two-input fallback rule base below is reserved for **Sentiment-Unreliable** sessions, specifically sessions where negativity cannot be resolved reliably after the available text or no-text path is attempted. This includes sessions whose extracted text is at least half unresolved vocabulary (study-defined unreliability trigger: session OOV ratio $\geq 50\%$) and sessions where the no-text VLM path cannot produce a stable item label because `AccessibilityService.takeScreenshot` is unavailable, screenshot capture fails, or VLM inference fails. The $\geq 50\%$ value is not treated as a validated linguistic cutoff. It is a mathematically derived majority-representativeness screen: once unresolved tokens reach half of the valid tokens in a session, the recognized tokens no longer form the majority of the lexical content available to the text path. This is therefore used as a practical safeguard for noisy code-mixed text with spelling variation, lexical borrowing, and other textual variation rather than as a claimed language norm (Mohammed & Prasad, 2023; Perera & Caldera, 2024; Nazir et al., 2026; Khan et al., 2025). In those cases, the system degrades to a two-input fallback rule base (Video Dwell Time + Session Duration):

Table 34. Fallback 9-Rule Base (Sentiment-Unreliable Sessions)

No.	Video Dwell Time	Session Duration	Risk Level
1	Low	Low	Safe
2	Low	Medium	Safe
3	Low	High	Warning
4	Medium	Low	Safe
5	Medium	Medium	Warning

No.	Video Dwell Time	Session Duration	Risk Level
6	Medium	High	Warning
7	High	Low	Warning

Table 7. Fallback 9-Rule Base (Sentiment-Unreliable Sessions) (continued)

No.	Video Dwell Time	Session Duration	Risk Level
8	High	Medium	Critical
9	High	High	Critical

In Week 1 and in any participant without locked personalization, both fallback inputs use the global memberships reported in Appendix C. Once Week 2 personalization is locked, the fallback engine retains the same 9-rule consequents and RiskScore cutoffs but replaces only the Session Duration memberships with participant-specific Week 1 duration quantiles; Video Dwell Time remains fixed. This is because the fallback mode reuses the same live prompt engine’s duration memberships while dropping only the sentiment input.

RiskScore cutoffs used across both full and fallback modes (equal-width reporting categories on the normalized 0-100 scale):

- Safe: 0.00 to 33.32
- Warning: 33.33 to 66.66
- Critical: 66.67 to 100

Prompt mapping (study-specified defaults derived from the reporting bands):

- Level 1: RiskScore 33.33 to 49.99
- Level 2: RiskScore 50.00 to 66.66
- Level 3: RiskScore 66.67 to 100 (disabled in fallback mode)
- The Warning band is split at its midpoint 50.00 so lower- and upper-Warning states receive different prompt intensities without introducing extra off-band cutoffs; Level 3 is reserved for the Critical band.
- Cooldown: fixed 15-minute prototype cooldown between prompts, intentionally reusing the same literature-bounded 15-minute live gate as the minimum re-

eligibility interval so the system does not introduce a second unsupported minute constant; a repeat prompt therefore requires another prolonged interval of continued use. Recent JITAI work supports managing timing, receptivity, and dose, but not this exact interval (Teepe et al., 2021; Wang et al., 2023; Fiedler et al., 2024; Hsu et al., 2025; van Genugten et al., 2025)

APPENDIX C

Core Algorithmic Logic and Reproducibility Specifications

The field-tested artifact separates fixed analytic priors from Week 2 live-prompt personalization. Fixed global memberships are retained for week-level DSI computation across both study weeks so baseline and intervention summaries remain comparable. A separate live prompt engine is then allowed to replace only the Session Duration and NSD memberships after Week 1; Video Dwell Time remains fixed.

Fixed Analytic Priors (used for week-level DSI across both study weeks and as the Week 1 live default): Where

$$\text{tri}(a, b, c)$$

is the standard triangular membership function, the study uses the following fixed memberships:

$$\text{Low}_{\text{Dwell}} = \text{tri}(0, 0, 5)$$

$$\text{Medium}_{\text{Dwell}} = \text{tri}(4, 12, 20)$$

$$\text{High}_{\text{Dwell}} = \text{tri}(15, 30, 30)$$

$$\text{Low}_{\text{NSD}} = \text{tri}(0, 0, 33)$$

$$\text{Medium}_{\text{NSD}} = \text{tri}(17, 50, 83)$$

$$\text{High}_{\text{NSD}} = \text{tri}(67, 100, 100)$$

$$\text{Low}_{\text{Duration}} = \text{tri}(0, 0, 10)$$

$$\text{Medium}_{\text{Duration}} = \text{tri}(8, 15, 20)$$

$$\text{High}_{\text{Duration}} = \text{tri}(15, 40, 40)$$

NSD values are clipped to $[0, 100]$. These fixed memberships govern week-level analytic DSI computation across both study weeks. In Week 1 they also act as the live default priors before any participant-specific personalization is available.

Where recent studies provide direct numeric anchors, only those anchors are treated as literature-bounded values: 5 s and 30 s for dwell time, 10 min, 15 min, and 40 min for session duration, and the normalized NSD/output partitions derived mathematically from the 0-100 scales. The remaining interior overlap points in the fixed priors are analytic

smoothing values used to preserve overlap between adjacent triangular sets and prevent the engine from collapsing into hard step thresholds. They are therefore treated as transition parameters and included in the sensitivity analysis rather than defended as standalone validated behavioral cutoffs.

Week 2 Prompt Personalization (implemented in the live prompt engine only): For participant u , let the Week 1 sentiment-reliable-session values for personalized variable X (Session Duration or NSD) be $x_{u,1}, x_{u,2}, \dots, x_{u,n_u}$. The prompt engine computes ordered quantiles:

$$Q_{25,u,X}, Q_{50,u,X}, Q_{75,u,X}, Q_{95,u,X}$$

The personalized Week 2 memberships are then:

$$\text{Low}_{X,\text{prompt},u} = \text{tri}(0, 0, Q_{50,u,X})$$

$$\text{Medium}_{X,\text{prompt},u} = \text{tri}(Q_{25,u,X}, Q_{50,u,X}, Q_{75,u,X})$$

$$\text{High}_{X,\text{prompt},u} = \text{tri}(Q_{50,u,X}, Q_{75,u,X}, Q_{95,u,X})$$

These quantiles are not treated as validated behavioral cutoffs. They are nonparametric order-statistic anchors chosen to map a three-term Low/Medium/High partition onto each participant's Week 1 distribution with the smallest transparent summary set: $\frac{Q_{25}}{Q_{75}}$ summarize the interquartile core, while Q_{95} caps the High set without letting one raw maximum set the upper endpoint. Video Dwell Time remains fixed to the global priors above. If $n_u < 10$ sentiment-reliable Week 1 sessions, the default priors are retained for that participant rather than locking personalized bounds. This 10-session minimum remains a numerical sufficiency safeguard rather than a validated population cutoff: below $n_u = 10$, the empirical Q_{95} is governed too heavily by the single largest observation or a nearly maximal order statistic, so the personalized High endpoint becomes too sensitive to one extreme session. Recent personalized intervention-criteria work supports deriving live criteria from a first-week baseline, but recent JITAI reviews also note that empirical decision rules and points remain underdeveloped, which is why the fielded artifact prefers participant-specific Week 2 criteria only when enough baseline data exist (Ikegaya et al., 2025; Hsu et al., 2025; van

Genugten et al., 2025; Elmer et al., 2025). Because fallback prompting uses the same live prompt engine, personalized Session Duration bounds also carry into fallback mode when personalization is available; only the rule base and prompt cap differ.

Risk Score Category Cutoffs (0-100 scale):

- Safe: $0 \leq \text{RiskScore} < 33.33$
- Warning: $33.33 \leq \text{RiskScore} < 66.67$
- Critical: $66.67 \leq \text{RiskScore} \leq 100$

Output Centers Used in CoG Aggregation:

- Safe center: 16.67
- Warning center: 50
- Critical center: 83.33

These fixed centers are the band midpoints of the equal-width thirds above, so they are mathematically derived ordinal anchors rather than empirical risk probabilities. This removes arbitrary-looking output constants while preserving smooth interpolation across mixed-rule activations before the final category label is assigned.

Intervention Mapping and Cooldown:

- No intervention when $\text{RiskScore} < 33.33$
- Level 1 (Awareness Notification): $33.33 \leq \text{RiskScore} < 50$
- Level 2 (Pause Prompt): $50 \leq \text{RiskScore} < 66.67$
- Level 3 (Breathing Exercise): $66.67 \leq \text{RiskScore} \leq 100$
- The Warning band is split at its midpoint 50.00 so lower- and upper-Warning states receive different prompt intensities without introducing extra off-band cutoffs; Level 3 is reserved for the Critical band.
- Fixed cooldown: minimum 15-minute prototype interval between any two prompts, intentionally reusing the same literature-bounded 15-minute live gate as

the minimum re-eligibility interval so the system does not introduce a second unsupported minute constant; a repeat prompt therefore requires another prolonged interval of continued use. Recent JITAI work supports managing timing, receptivity, and dose, but not this exact interval (Teepe et al., 2021; Wang et al., 2023; Fiedler et al., 2024; Hsu et al., 2025; van Genugten et al., 2025)

When the 2-input behavioral fallback mode is active (Sentiment-Unreliable), Level 3 is disabled by safety lock and the maximum permitted prompt is Level 2.

APPENDIX D

Short-Form Doomscrolling Scale Administration Sheet

The study uses the short-form 4-item **Doomscrolling Scale**, corresponding to Items 1, 2, 10, and 12 of the original 15-item instrument developed by Sharma et al. (2022) and later confirmed in short-form use by Satici et al. (2023). To align self-report with each study week, the instrument is administered with a 7-day recall instruction stem. This modification narrows the recall window only; it does not create a separately validated weekly state instrument.

Participant Instructions: “In the past week, please indicate how strongly you agree with each statement.”

Response Format:

- 1 = Strongly disagree
- 2 = Disagree
- 3 = Somewhat disagree
- 4 = Neither agree nor disagree
- 5 = Somewhat agree
- 6 = Agree
- 7 = Strongly agree

Items:

1. I feel an urge to seek bad news on social media, more and more often.
2. I lose track of time when I read bad news on social media.
3. I find myself continuously browsing negative news.
4. I feel like I am addicted to negative news.

Scoring Note: Item responses are summed or averaged to produce a single doomscrolling score, with higher values indicating greater self-reported doomscrolling during the recalled 7-day period. The 4-item short form contains no reverse-coded items.